

# Distributed Systems

殷亚凤

Email: [yafeng@nju.edu.cn](mailto:yafeng@nju.edu.cn)

Homepage: <http://cs.nju.edu.cn/yafeng/>

Room 301, Building of Computer Science and Technology



# Content

- Lesson: 10% (13+1 times)
- Homework: 30% (2 times: Survey, Project)
- Exam (closed-book): 60%
- Lecture: Cloud computing, Big data, IoT, OpenStack (3 times)

# Lesson

- Introduction
- System Architecture
- Processes
- **Communication**
- **Synchronization**
- **Consistency** & Replication
- **Fault Tolerance**
- Distributed Object
- Distributed File Systems
- Distributed Web-based Systems


- **Textbook**

Andrew S. Tanenbaum and Maarten van Steen, *Distributed Systems: Principles and Paradigms*, 2nd edition, Prentice Hall, 2007.




- **Reference**


George Coulouris etc., *Distributed Systems: Concepts and Design*, 5th edition, Addison-Wesley, 2005.




我们是否处于  
“第三次工业革命”的前夜？



或者说：  
又一次“工业革命”已经开始！



科技革命





# Everything Online

高度互联: anytime, anywhere, any way!

**物联网:**

Everything can be connected.



**大数据:**

A **revolution** that will transform how we live, work, and think!



**人工智能:**

Simulation and extension of human intelligence.



# 什么是背后的推动力？

计算机科学在新一轮变革中必将承担以前数学与自然科学所承担的使命。

分布式系统是构成当代信息系统的重要基础架构。

# Introduction

Distributed Systems [1]

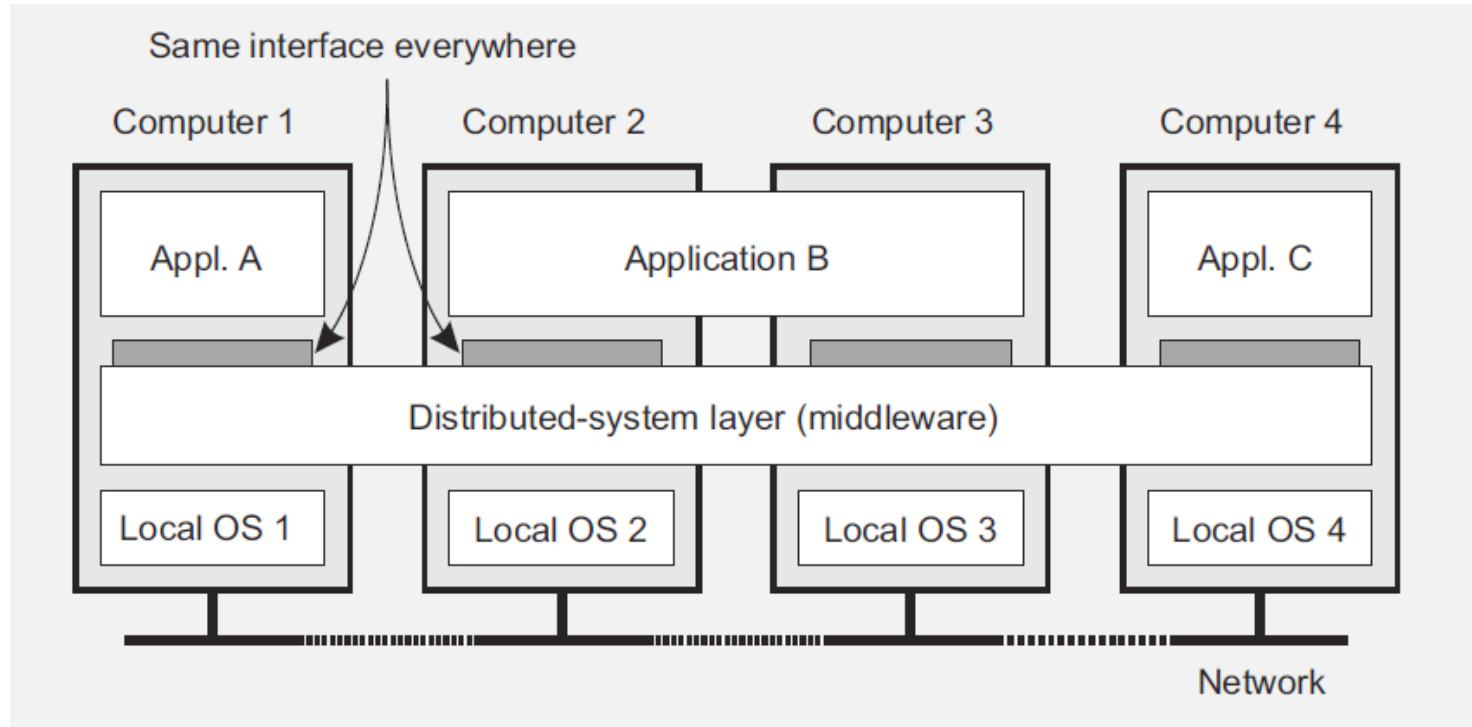


# This lesson

- **Definition of distributed systems**
- **Goals of Distributed Systems**
- **Types of Distributed Systems**

# Distributed Systems: Definition

- A distributed system is a collection of *autonomous computing elements* that appears to its users as a *single coherent system*.



# Examples

- A network of **workstations** allocated to users
- A pool of **processors** in the machine room allocated dynamically
- A single file system (all users **access files** with the same path name)
- **User command** executed in the best place (user workstation, a workstation belonging to someone else, or on an unassigned processor in the machine room)

# Technology advances

- Networking
- Processors
- Memory
- Storage
- Protocol

# Why Distributed?

Economics	Microprocessors offer a better price/performance than mainframes
Speed	A distributed system may have more total computing power than a mainframe
Inherent distribution	Some applications involve spatially separated machines
Reliability	If one machine crashes, the system as a whole can still survive
Incremental growth	Computing power can be added in small increments

People are distributed

# Goals of Distributed Systems

- Making resources available
- Distribution transparency
- Openness
- Scalability

# Transparency in a Distributed System

Transparency	Description
<b>Access</b>	Hide differences in data representation and how a resource is accessed
<b>Location</b>	Hide where a resource is located
<b>Migration</b>	Hide that a resource may move to another location
<b>Relocation</b>	Hide that a resource may be moved to another location while in use
<b>Replication</b>	Hide that a resource may be shared by several competitive users
<b>Concurrency</b>	Hide that a resource may be shared by several competitive users
<b>Failure</b>	Hide the failure and recovery of a resource

Distribution transparency is a nice goal, but achieving it is a different story.



# Degree of Transparency

- Aiming at **full distribution transparency** may be too much difficult:
  - Users may be located in different **continents**
  - Completely **hiding failures** of networks and nodes is (theoretically and practically) impossible
    - You cannot distinguish a slow computer from a failing one
    - You can never be sure that a server actually performed an operation before a crash
  - Full transparency will cost performance, **exposing distribution** of the system
    - Keeping Web caches exactly up-to-date with the master
    - Immediately flushing write operations to disk for fault tolerance

# Openness of Distributed Systems

- Be able to **interact** with services from other open systems, irrespective of the underlying environment:
  - Systems should conform to well-defined **interfaces**
  - Systems should support **portability** of applications
  - Systems should easily **interoperate**
- At least make the distributed system **independent from heterogeneity** of the underlying environment.

# Policies Versus Mechanisms

- Requires support for different **policies**:
  - What level of **consistency** do we require for client-cached data?
  - Which **operations** do we allow downloaded code to perform?
  - Which **QoS requirements** do we adjust in the face of varying bandwidth?
  - What **level of secrecy** do we require for communication?
- Ideally, a distributed system provides only **mechanisms**:
  - Allow (dynamic) **setting** of caching policies
  - Support different levels of **trust** for mobile code
  - Provide adjustable **QoS parameters** per data stream
  - Offer different **encryption** algorithms

# Scalability

- At least three components:
  - Number of users and/or processes (**size** scalability)
  - Maximum distance between nodes (**geographical** scalability)
  - Number of administrative domains (**administrative** scalability)
- Most systems account only, to a certain extent, for **size scalability**. The (non)solution: powerful servers. Today, the challenge lies in **geographical** and **administrative** scalability.

# Types of Distributed Systems

- Distributed computing systems
- Distributed information systems
- Distributed pervasive systems

# Distributed Computing Systems


- **Cluster Computing**

- Essentially a group of high-end systems connected through a LAN:
  - **Homogeneous**: same OS, near-identical hardware
  - Single managing node

- **Grid Computing**

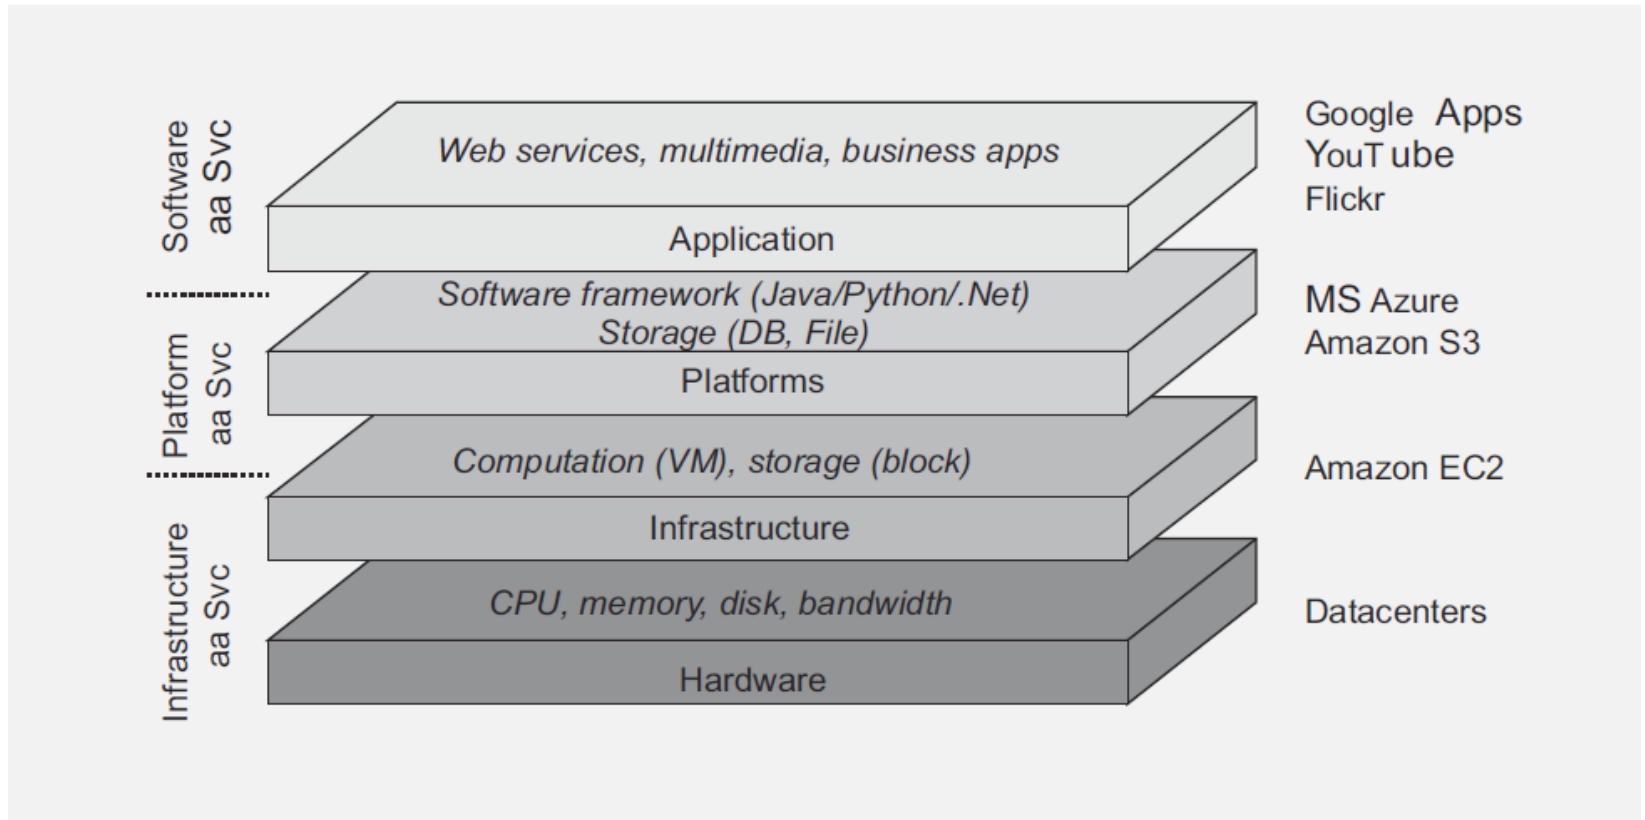
- The next step: lots of nodes from everywhere:
  - **Heterogeneous**
  - Dispersed across several organizations
  - Can easily span a wide-area network

# Distributed Computing Systems: Clouds





# Cloud Platform

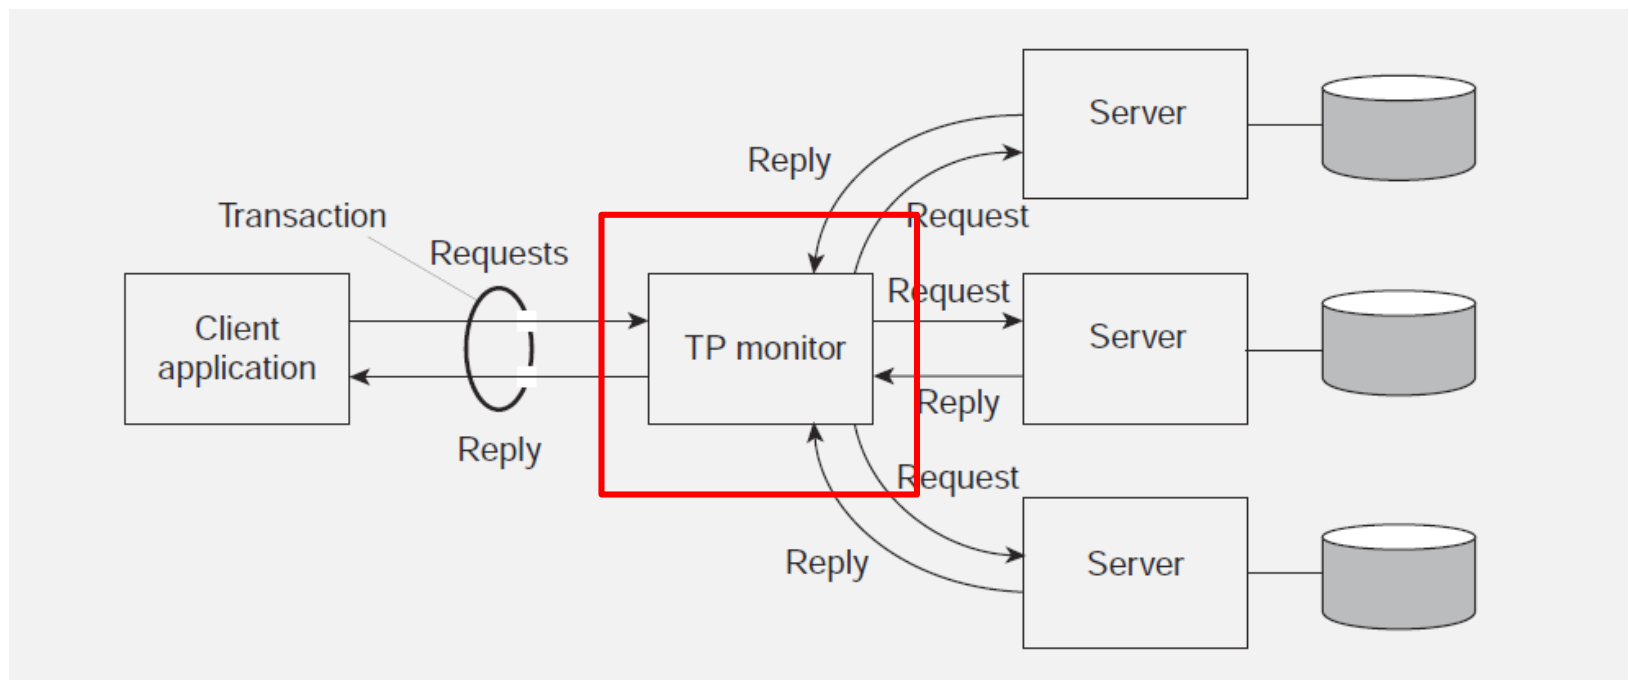


# Distributed Information Systems

- The vast amount of distributed systems in use today are forms of traditional information systems, that now integrate legacy systems.
  - Example: [Transaction processing systems](#).
- A transaction is a collection of operations on the state of an object (database, object composition, etc.) that satisfies the following properties (ACID)
  - [Atomicity](#)
  - [Consistency](#)
  - [Isolation](#)
  - [Durability](#)

# Transaction Processing Monitor

- In many cases, the data involved in a transaction is distributed across several servers. A **TP Monitor** is responsible for coordinating the execution of a transaction.




# Distributed Pervasive Systems

- Emerging **next-generation of distributed systems** in which nodes are **small, mobile**, and often **embedded in a larger system**, characterized by the fact that the system naturally blends into the user's environment.
  - **Ubiquitous** computing systems
  - **Mobile** computing systems
  - **Sensor** (and actuator) **networks**


# Some Typical Examples

- Internet
- Intranet
- Mobile environment
- Web

# Internet

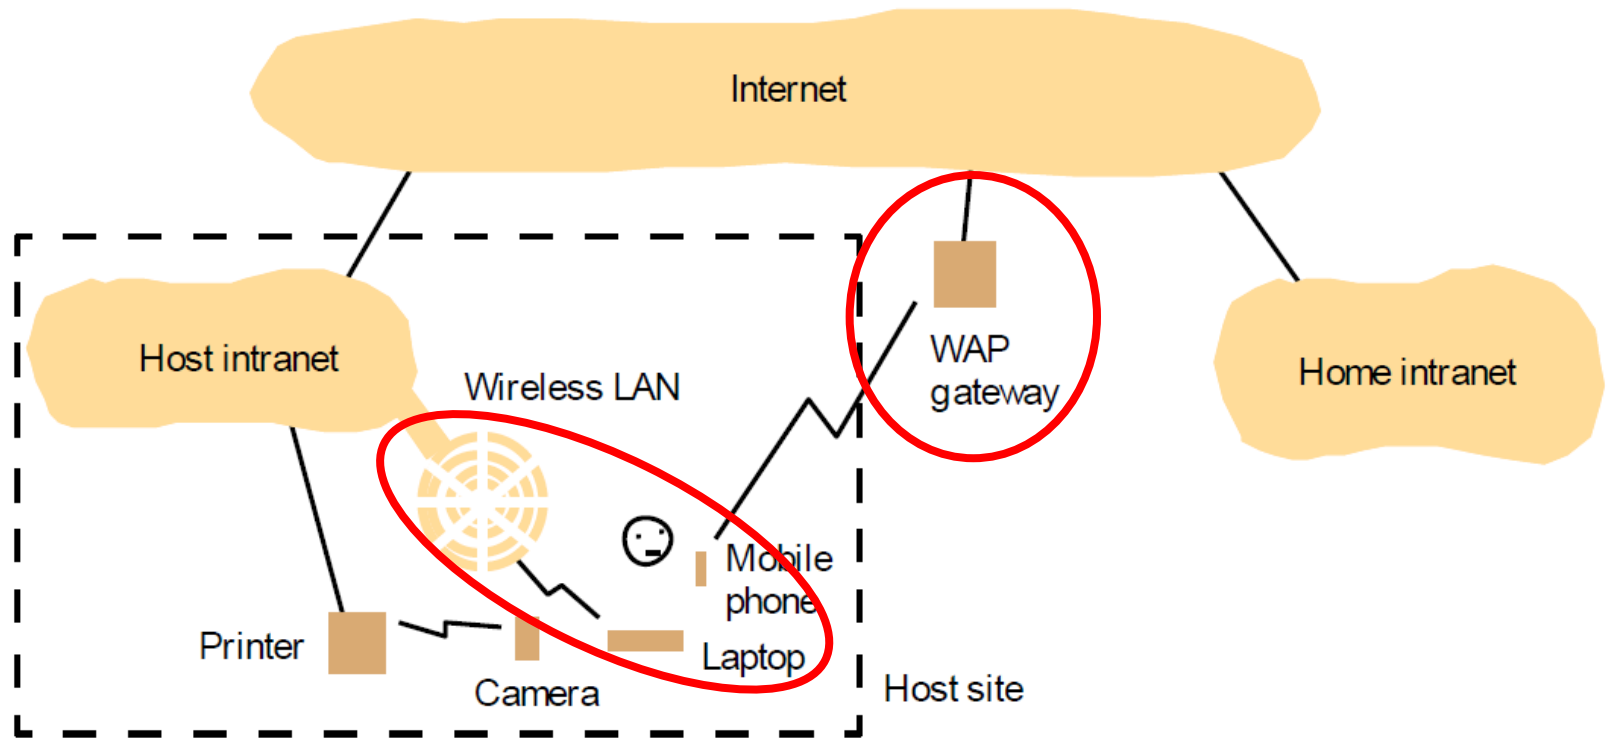


# Intranet





# Mobile Environment



# Web

