

Processes

Distributed Systems [3]

殷亚凤

Email: yafeng@nju.edu.cn

Homepage: <http://cs.nju.edu.cn/yafeng/>

Room 301, Building of Computer Science and Technology

Review

- **Architecture Styles:** Layered, object-oriented, event-based, shared data spaces-based
- **System Architecture :** Centralized, Decentralized, Hybrid
- **Middleware**
- **Self-managing Distributed Systems**

This lesson

- **Process**
- **Thread**
- **Client – Server**
- **Code Migration**

What is a Process?


- Process: An **execution stream** in the context of a **process state** (a program in execution)
- **Execution stream**
 - Stream of executing instructions
 - Running piece of code
 - Sequential sequence of instructions
 - “thread of control”
- **Process state**
 - Everything that the running code can affect or be affected by

Processes vs. Programs

- A process is different than a program
 - Program: Static code and static data
 - Process: Dynamic instance of code and data
- No one-to-one mapping between programs and processes
 - One process can execute multiple programs
 - One program can invoke multiple processes

Processes

- Each process is in one of three modes:
 - **Running**: On the **CPU** (only one on a uniprocessor)
 - **Ready**: Waiting for the **CPU**
 - **Blocked** (or asleep): Waiting for I/O or synchronization to complete



Low Performance

- **Resource management**

- When **creating** a new process, assign **address space**, **copy data**

- **Scheduling**

- Context **switch**
 - Process Context: **CPU** context and storage context

- **Cooperation**

- **IPC**, interprocess communication
 - Shared memory

Introduction to Threads

- **Thread**: A **minimal software processor** in whose context a series of instructions can be executed.
- **Saving** a thread context implies stopping the current execution and saving **all the data needed to** continue the execution at a later stage.

Process and thread : Context Switching


- **Processor context:** The minimal collection of values stored in the **registers of a processor** used for the execution of a series of **instructions** (e.g., stack pointer, addressing registers, program counter).
- **Thread context:** The minimal collection of values stored in **registers and memory**, used for the execution of a series of **instructions** (i.e., processor context, state).
- **Process context:** The minimal collection of values stored in **registers and memory**, used for the execution of a **thread** (i.e., thread context, but now also at least MMU register values).

Context Switching


- **Threads** share the same address space. Thread context switching can be done entirely **independent of the operating system**.
- **Process** switching is generally more expensive as it involves getting the **OS in the loop**, i.e., to the kernel.
- **Creating and destroying** threads is much cheaper than doing so for processes.

Thread and Process

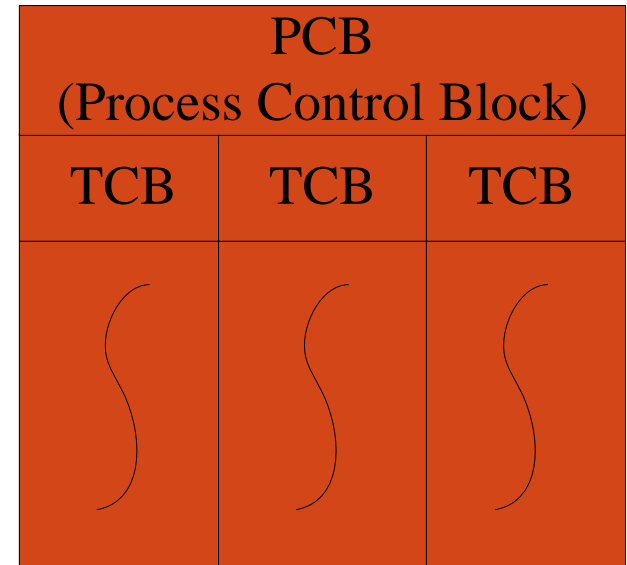
Single Thread
Process



Multi Thread
Process



Multi Thread
Process



Tread System

Operation System


User-Level Thread

- All the threads are created in user processes' address spaces.
- **Advantage**
 - All operations can be completely handled within a single process \Rightarrow implementations can be extremely efficient.
- **Disadvantage**
 - Difficult to get the support from OS, block
 - All services provided by the kernel are done on behalf of the process in which a thread resides \Rightarrow if the kernel decides to block a thread, the entire process will be blocked.

Threads and Operating Systems

- Have the **kernel** contain the implementation of a thread package. This means that **all operations return as system calls**:
 - Operations that block a thread are no longer a problem: the **kernel schedules another available thread** within the same process.
 - Handling **external events** is simple: the kernel (which catches all events) schedules the thread associated with the event.
 - The problem is (or used to be) **the loss of efficiency** due to the fact that each thread operation requires a trap to the kernel.

Thread Implementation



Light weight processes, LWP

Threads and Distributed Systems

- Hiding network latencies:
 - Web browser scans an incoming **HTML page**, and finds that more files need to be fetched.
 - **Each file** is fetched by a separate thread, each **doing a (blocking) HTTP request**.
 - As files come in, the browser displays them.
- Multiple request-response calls to other machines(RPC)
 - A client does several calls at the same time, each one by a **different thread**.
 - It then **waits until all results** have been returned.
 - Note: if calls are to different servers, we may have a linear speed-up.

Multithreaded Clients

- Example, web browser

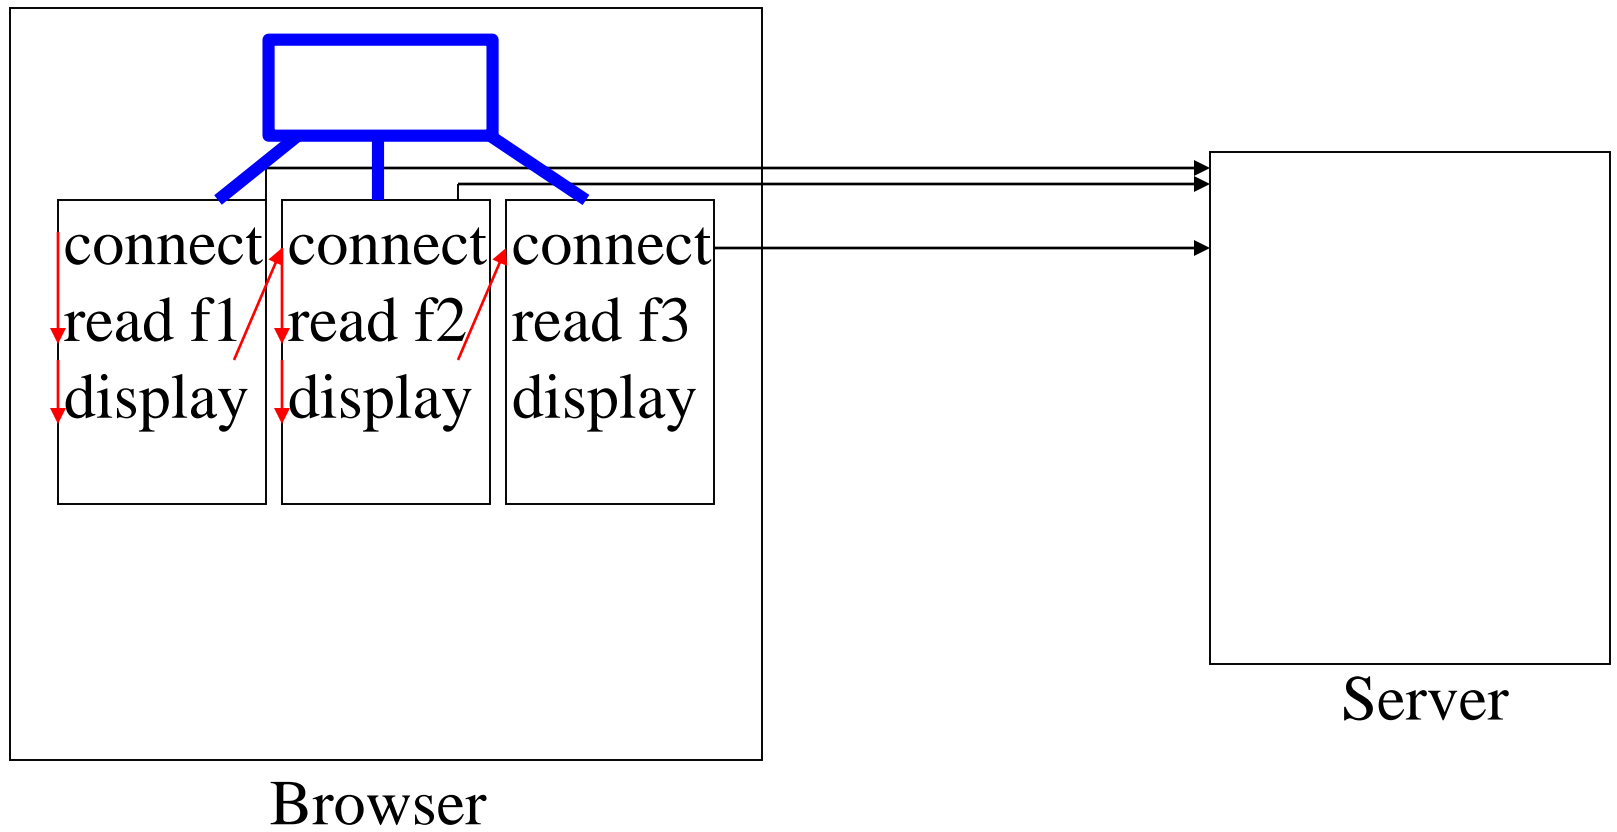
A web document \supset plain text, a collection of images.

To fetch a HTML file:

connect server, read a file1, display

connect server, read a file2, display

... ..

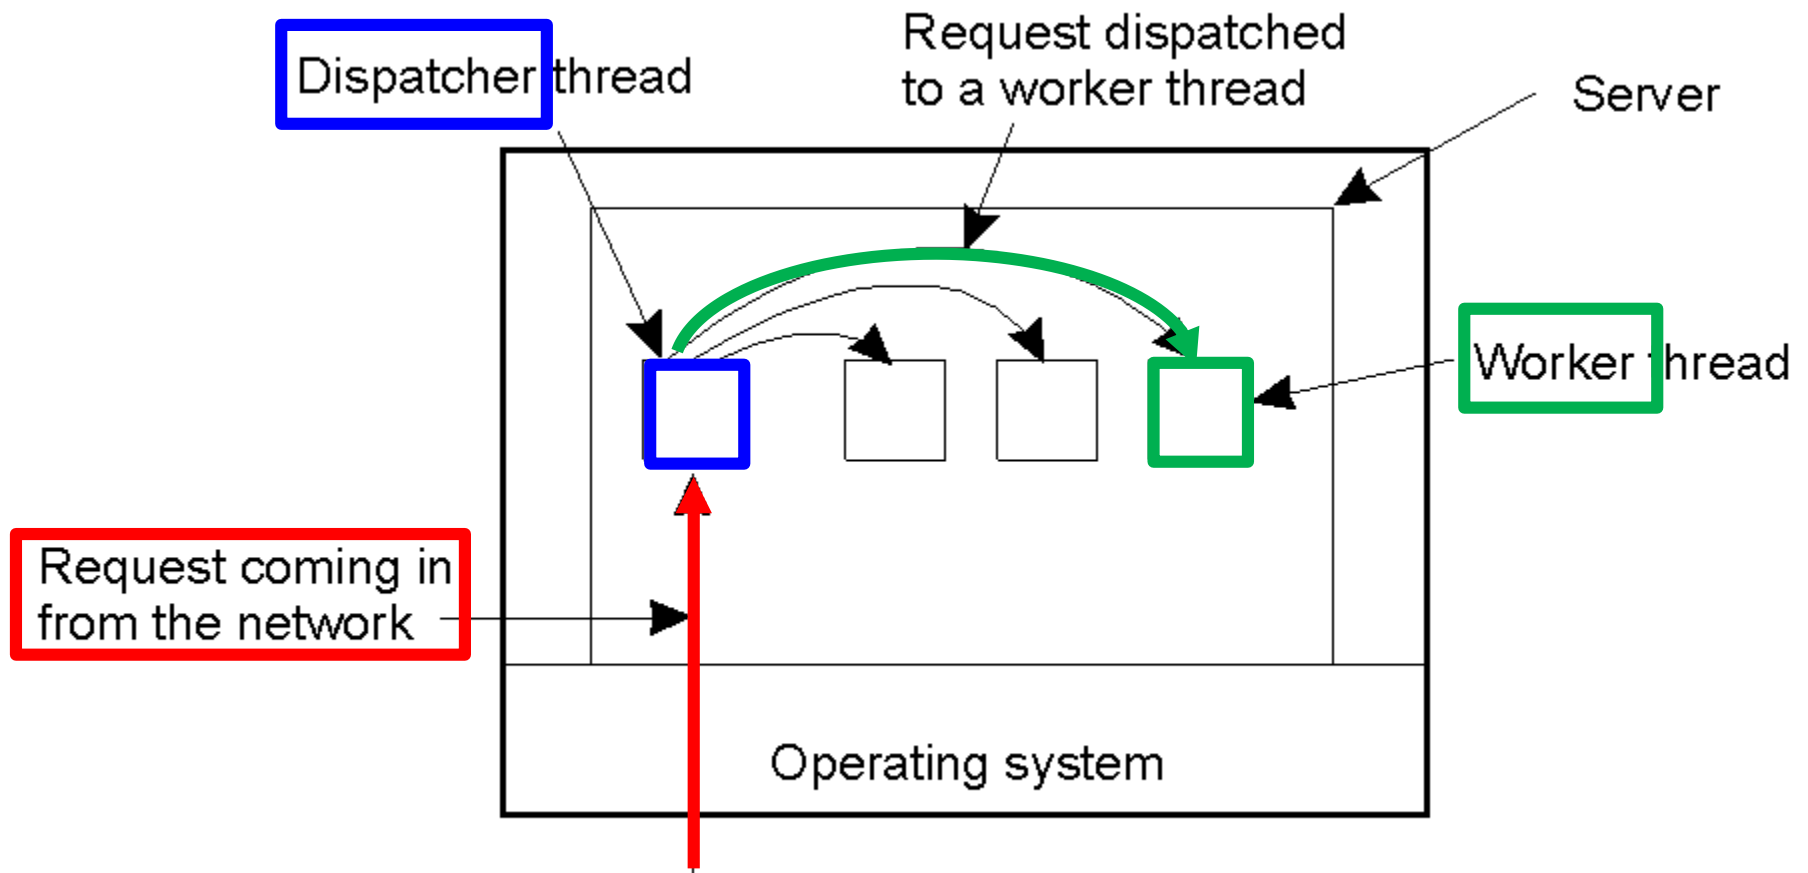


Threads and Distributed Systems

- **Improve performance:**
 - Starting a thread is much **cheaper** than starting a new process.
 - Having a single-threaded server prohibits simple scale-up to **a multiprocessor system**.
 - As with clients: **hide network latency** by reacting to next request while previous one is being replied.
- **Better structure**
 - Most servers have high I/O demands. Using simple, well-understood blocking calls **simplifies the overall structure**.
 - Multithreaded programs tend to be smaller and easier to understand due to **simplified flow of control**.


Multithreaded Servers

- A multithreaded server organized in a **dispatcher/worker** model.



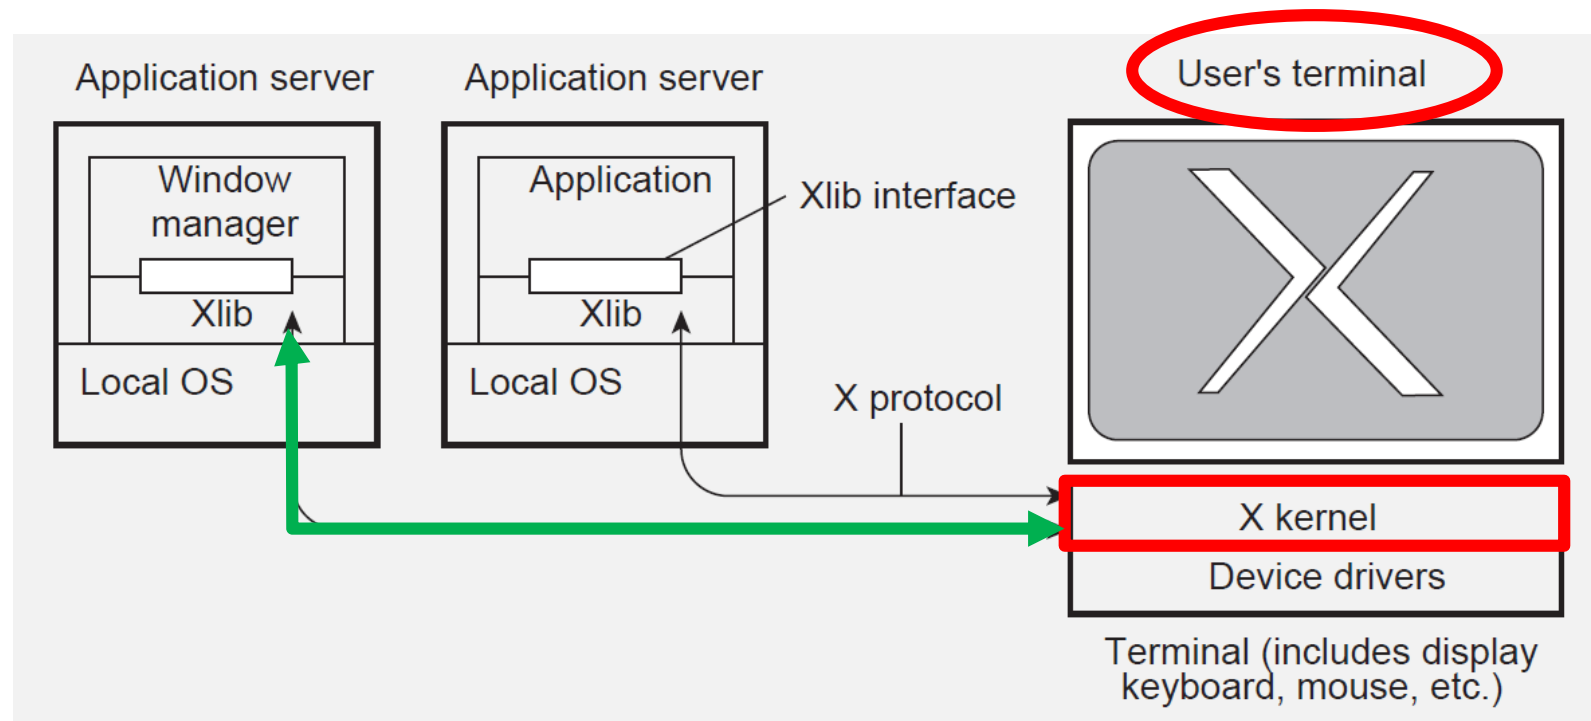
Multithreaded servers

- Example, file server



Clients: User Interfaces


- A major part of **client-side software** is focused on (graphical) **user interfaces**.



X Window

Client-Side Software

- Generally tailored for **distribution transparency**
 - **access** transparency: client-side stubs for RPCs
 - **location/migration** transparency: let client-side software keep track of actual location
 - **failure** transparency: can often be placed only at client (we're trying to mask server and communication failures).
 - **replication** transparency: multiple invocations handled by client stub



Servers: General organization

- Basic model:
 - A server is a process that **waits for incoming service** requests at a specific transport address. In practice, there is a one-to-one mapping between **a port and a service**.

ftp-data	20	File Transfer [Default Data]
ftp	21	File Transfer [Control]
telnet	23	Telnet
	24	any private mail system
smtp	25	Simple Mail Transfer
login	49	Login Host Protocol
sunrpc	111	SUN RPC (portmapper)
courier	530	Xerox RPC

Servers: General organization

- Type of servers:
 - **Superservers**: Servers that listen to **several ports**, i.e., provide several independent services. In practice, when a service request comes in, they start a **subprocess** to handle the request (UNIX inetd)
 - **Iterative vs. concurrent servers**: Iterative servers can handle only **one client at a time**, in contrast to **concurrent servers**

Out-of-band communication

- Is it possible to **interrupt** a server once it has accepted (or is in the process of accepting) a service request?
- Use a separate port for urgent data:
 - Server has a separate thread/process for urgent messages
 - Urgent message comes in \Rightarrow associated request is put on hold
 - Note: we require OS supports priority-based scheduling
- Use out-of-band communication facilities of the transport layer:
 - Example: TCP allows for urgent messages in same connection
 - Urgent messages can be caught using OS signaling techniques

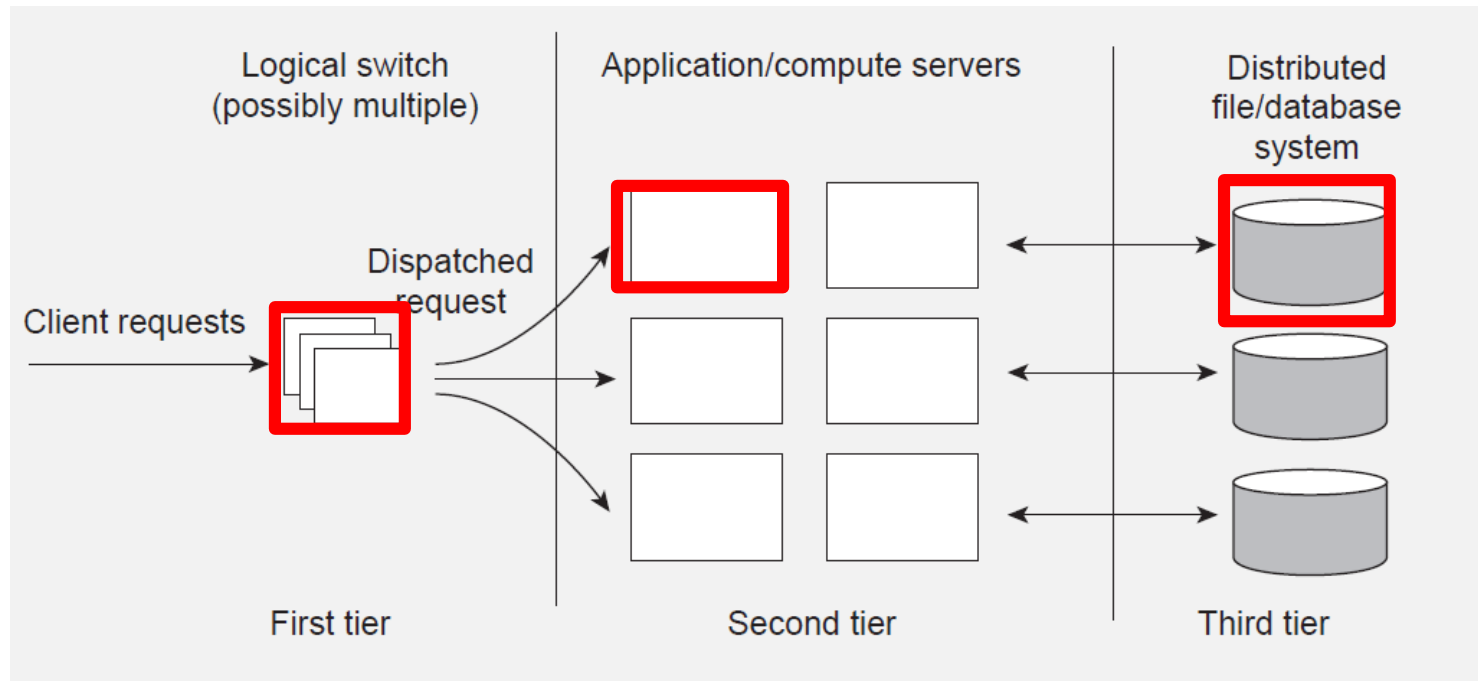
Stateless

- Never keep accurate information about the status of a client after having handled a request:
 - Don't record whether a file has been opened (simply close it again after access)
 - Don't promise to invalidate a client's cache
 - Don't keep track of your clients
- Consequences
 - Clients and servers are completely **independent**
 - State **inconsistencies** due to client or server crashes are **reduced**
 - Possible **loss of performance** because, e.g., a server cannot anticipate client behavior (think of prefetching file blocks)

Stateful

- Keeps track of the status of its clients:
 - **Record** that a file has been opened, so that prefetching can be done
 - Knows which data a client has cached, and allows **clients to keep local copies** of shared data
- The **performance** of stateful servers can be extremely **high**, provided clients are allowed to keep local copies. As it turns out, reliability is not a major problem.

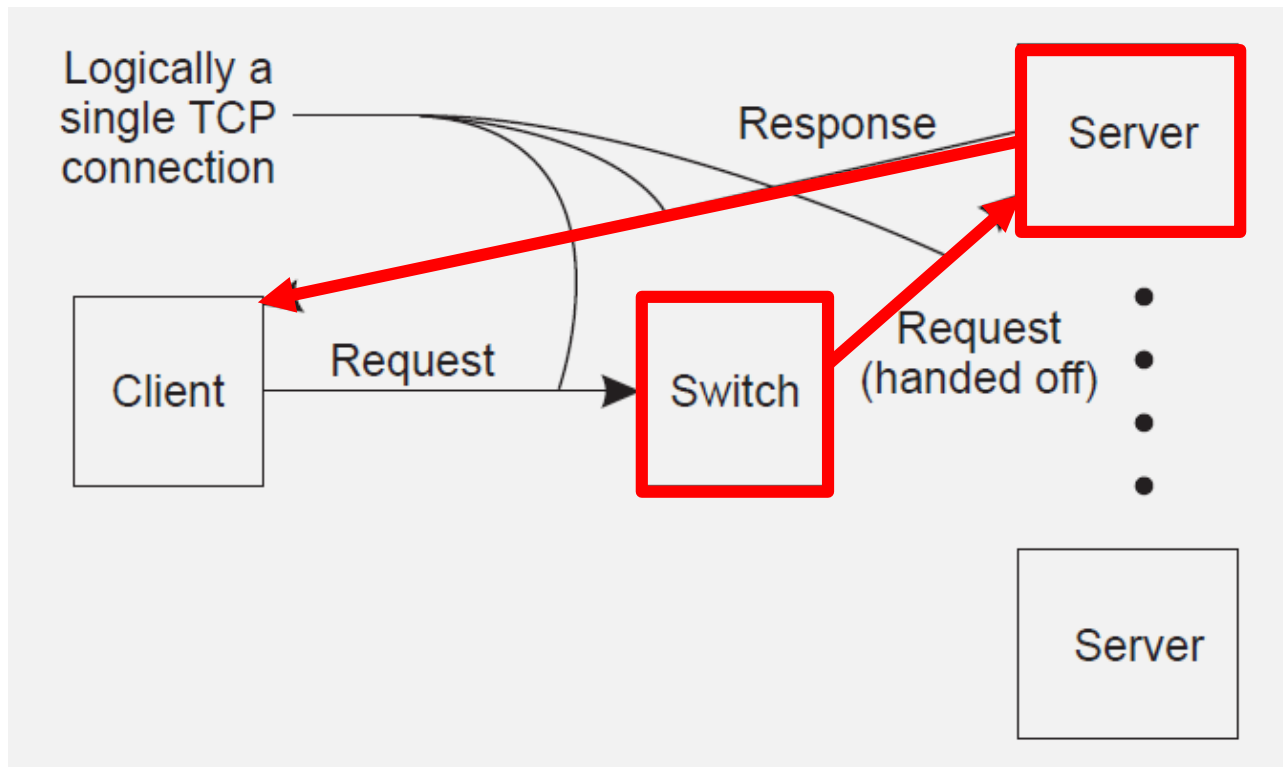
Server clusters: three different tiers



- Crucial element
 - The first tier is generally responsible for passing requests to an appropriate server.

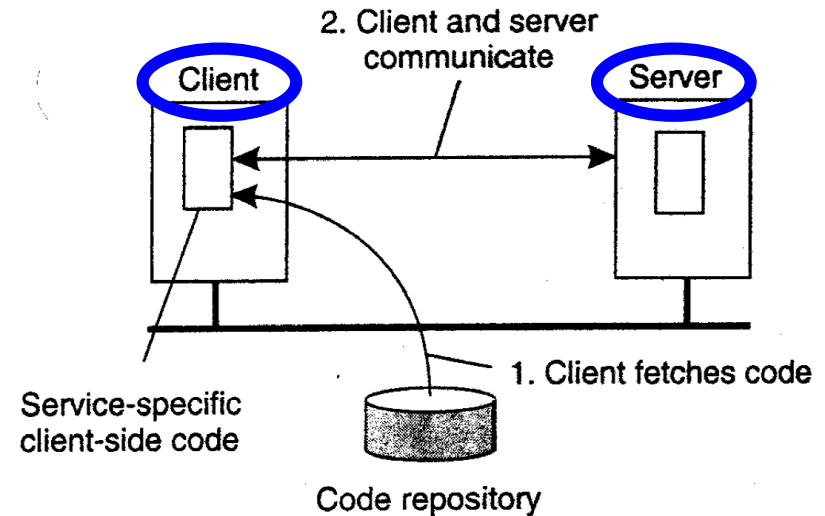
Request Handling

- The **first tier** handle all communication from/to the.
- The one popular one is **TCP-handoff**



Code Migration

- **Approaches** to code migration



- Migration and **local resources**
- Migration in **heterogeneous systems**

Migration model

- **Code segment**: contains the executing instructions
- **Resource segment**: contains the pointers pointing to the external resources
- **Execution segment**: contains executing states

Weak and strong mobility

Weak

- Move only **code and data segment** (and reboot execution):
 - Relatively simple, especially if code is portable
 - Distinguish code shipping (push) from code fetching (pull)

Strong

- Move component, including **execution state**
 - **Migration**: move entire object from one machine to the other
 - **Cloning**: start a clone, and set it in the same execution state

Managing local resources

- Object-to-resource binding
 - **By identifier**: the object requires a **specific instance** of a resource (e.g. a specific database)
 - **By value**: the object requires the **value** of a resource (e.g. the set of cache entries)
 - **By type**: the object requires that only a **type** of resource is available (e.g. a color monitor)

Managing local resources

- An object uses **local resources** that may or may not be available at the target site
- Resource types
 - **Fixed**: the resource cannot be migrated, such as local hardware
 - **Fastened**: the resource can, in principle, be migrated but only at high cost
 - **Unattached**: the resource can easily be moved along with the object (e.g. a cache)

Managing local resources (2/2)

	Unattached	Fastened	Fixed
ID	MV (or <u>GR</u>)	<u>GR</u> (or MV)	<u>GR</u>
Value	CP (or MV, <u>GR</u>)	<u>GR</u> (or CP)	<u>GR</u>
Type	<u>RB</u> (or MV, <u>GR</u>)	<u>RB</u> (or <u>GR</u> , CP)	<u>RB</u> (or <u>GR</u>)
<p><u>GR</u> = Establish global systemwide reference MV = Move the resource CP = Copy the value of the resource RB = Re-bind to a locally available resource</p>			

Migration in heterogeneous systems

- Main problem
 - The target machine **may not be suitable to execute the migrated code**
 - The definition of process/thread/processor context is highly dependent on **local hardware**, operating system and runtime system
- Make use of an **abstract machine** that is implemented on different platforms
 - **Interpreted languages**, effectively having their own VM
 - **Virtual VM** (as discussed previously)