

Fault Tolerance

Distributed Systems [8]

殷亚凤

Email: yafeng@nju.edu.cn

Homepage: <http://cs.nju.edu.cn/yafeng/>

Room 301, Building of Computer Science and Technology

Review

- Replication
- Data-Centric Consistency Models
- Client-Centric Consistency Models
- Replication Management
- Consistency Protocols

This Lesson

- Concepts about faults
- How to improve dependability
- Two-army problem
- Byzantine agreement problem

Basic Concepts

- Dependability Includes
 - Availability
 - Reliability
 - Safety
 - Maintainability

Dependability

- Availability

- The fraction of the time that a system meets its specification.
- The **probability that the system is operational** at a given time t .

- Reliability

- A measure of success with which a system conforms to some **authoritative** specification of its behavior.
- Probability that the system has **not experienced any failures** within a given time period.
- Typically used to describe systems that cannot be repaired or where the **continuous operation** of the system is critical.

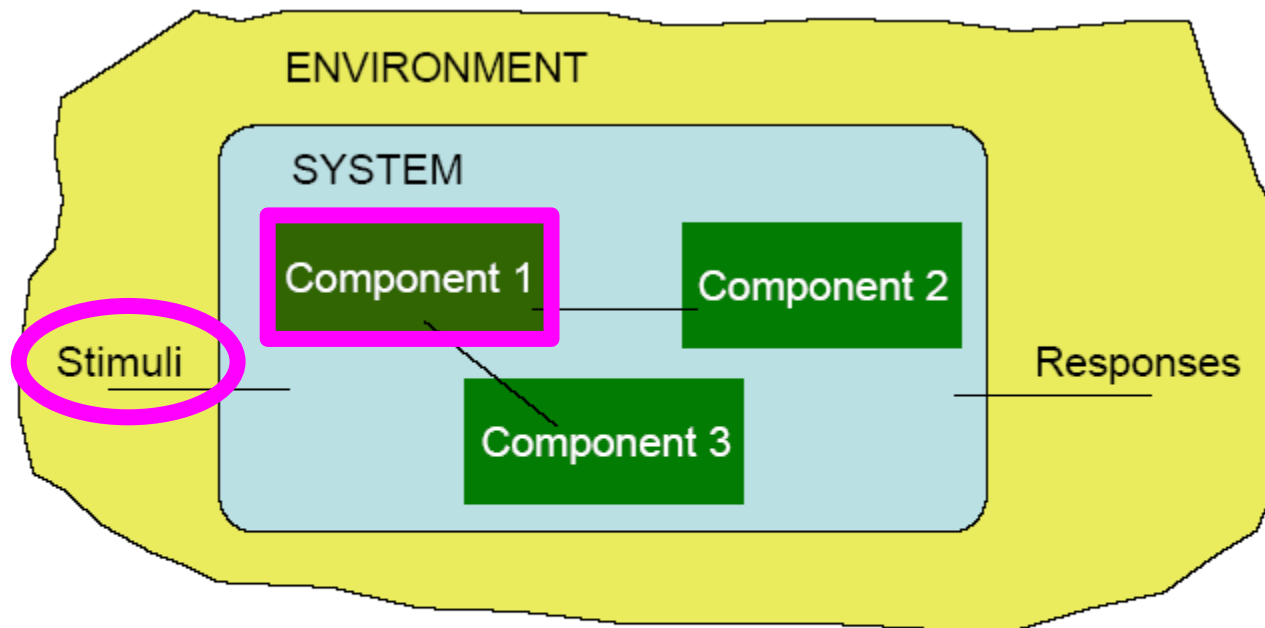
- Safety

- When the system temporarily fails to conform to its specification, **nothing catastrophic** occurs.

- Maintainability

- Measure of how easy it is to **repair** a system.

Basic System Concepts



External state

Internal state

Fundamental Definitions

- Failure

- A system is said to fail when it **cannot meet its promises**.

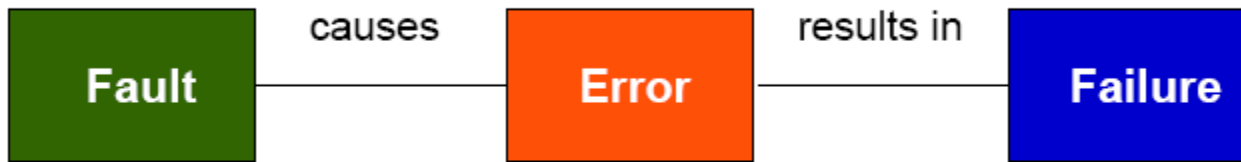
- Error

- The **part of the a system's state** that may lead to a failure.

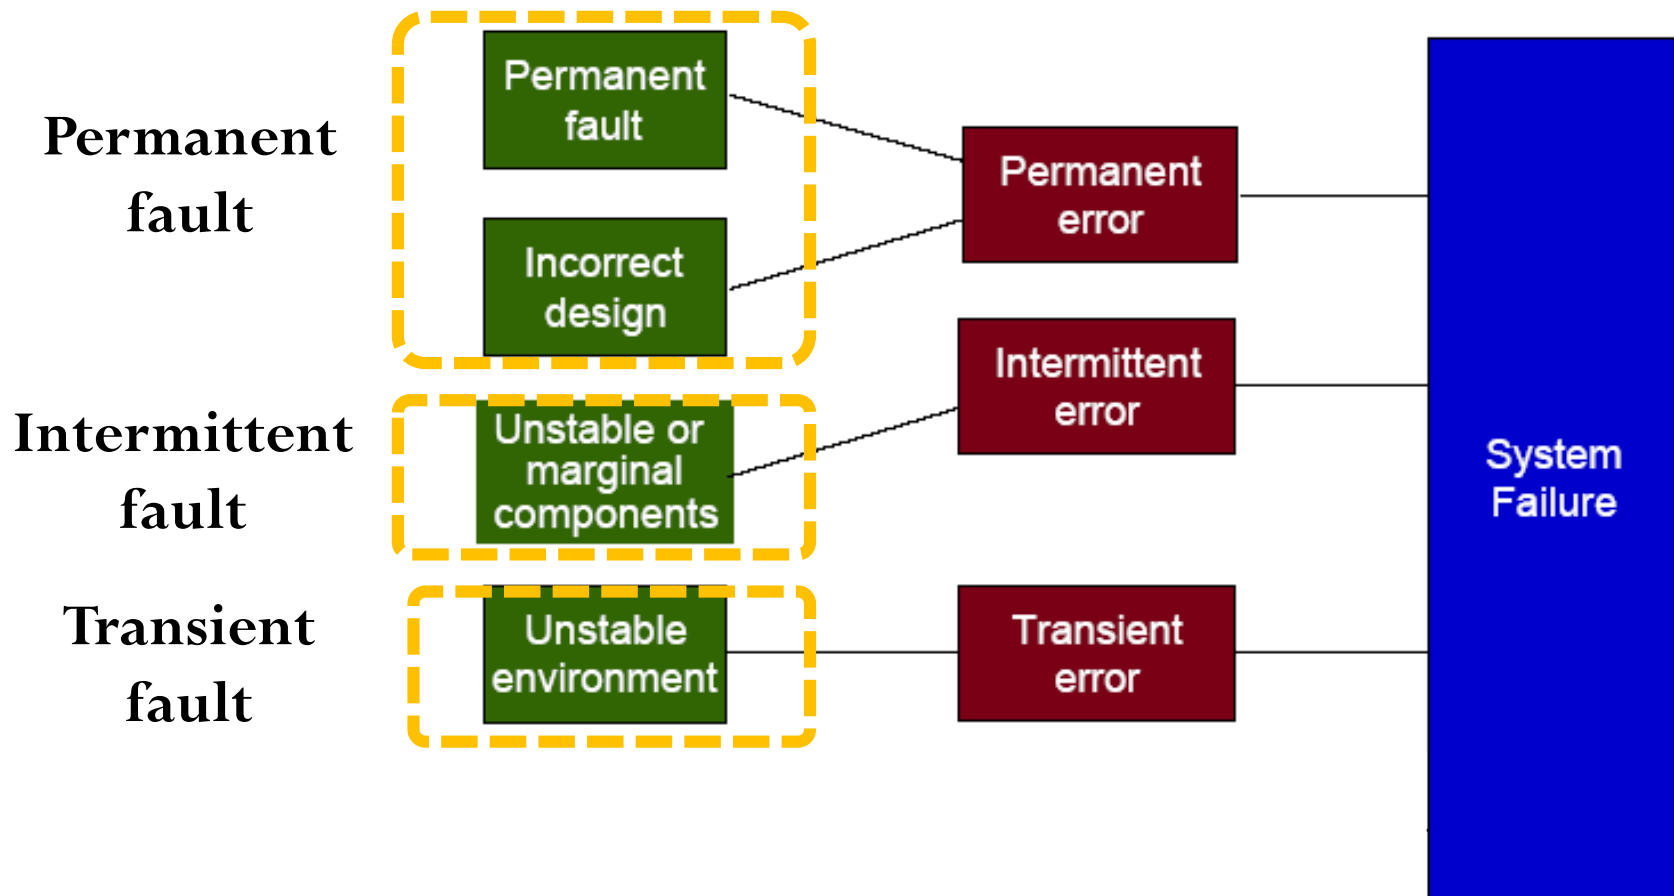
- Fault

- The **cause of an error** is called a fault.

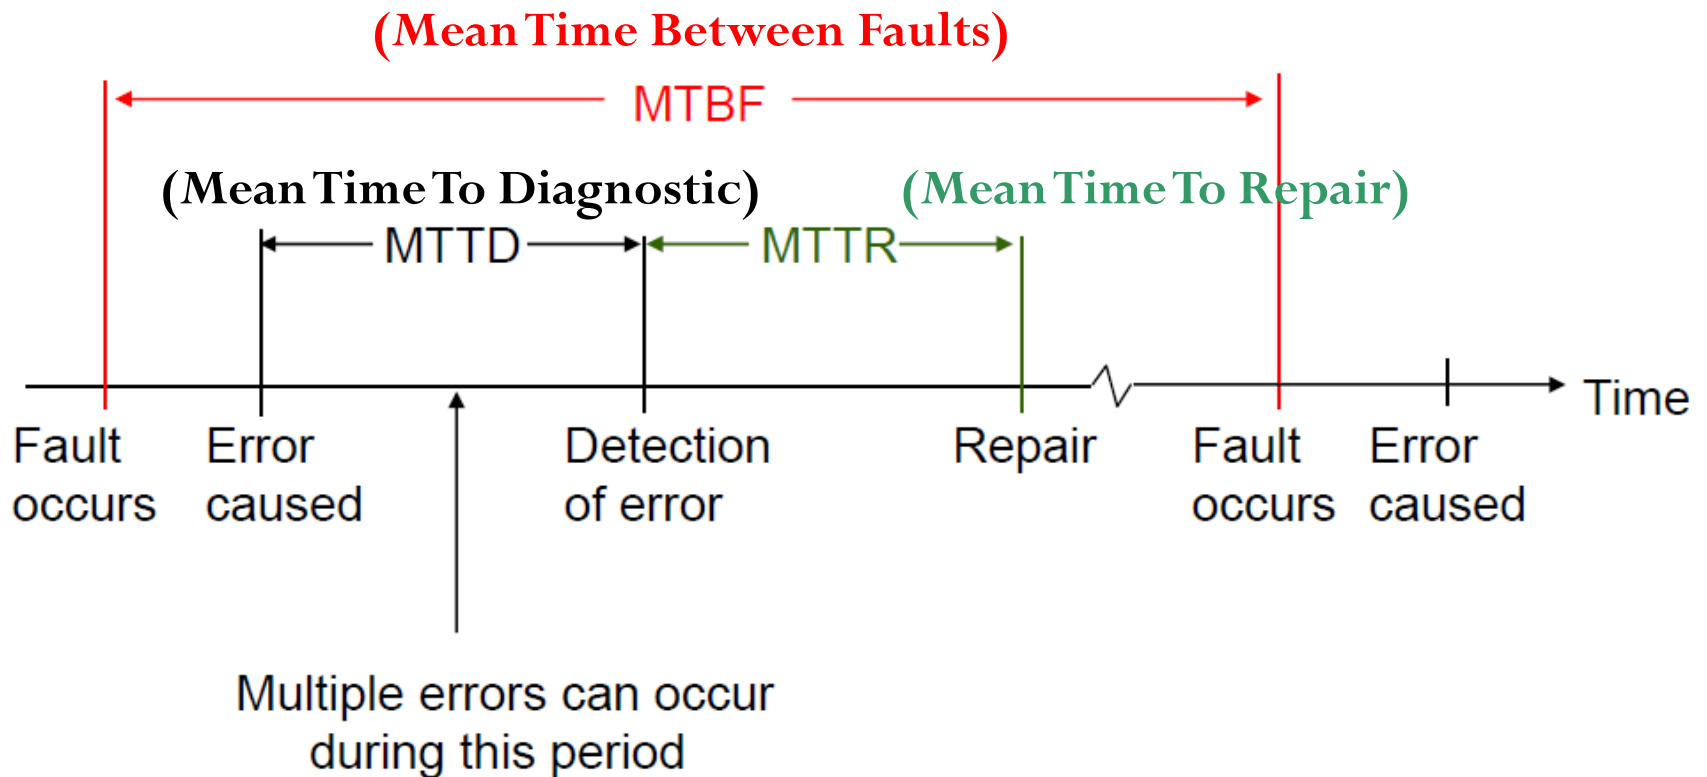
Faults to Failures



Fault Classification



Faults



Faults Models

- Different types of faults.

Type of failure	Description
Crash fault	A server halts , but is working correctly until it halts
Omission fault <i>Receive omission</i> <i>Send omission</i>	A server fails to respond to incoming requests A server fails to receive incoming messages A server fails to send messages
Timing fault	A server's response lies outside the specified time interval
Response fault <i>Value failure</i> <i>State transition failure</i>	The server's response is incorrect The value of the response is wrong The server deviates from the correct flow of control
Arbitrary fault	A server may produce arbitrary responses at arbitrary times

How to Improve Dependability

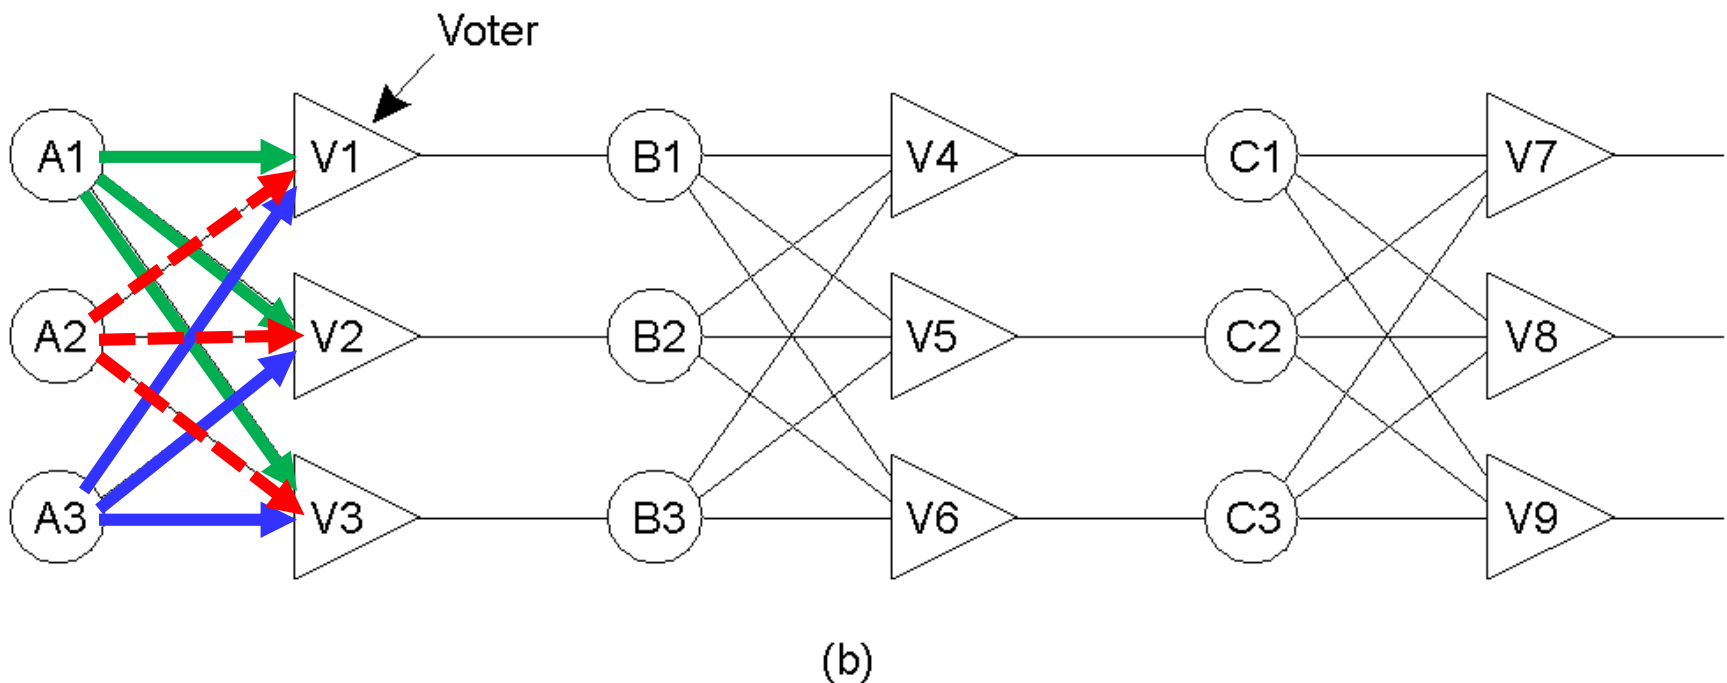

- Mask failures by redundancy
 - Information redundancy
 - E.g., **add extra bits** to detect and recover data transmission errors
 - Time redundancy
 - **Transactions**; e.g., when a transaction aborts re-execute it without adverse effects.
 - Physical redundancy
 - **Hardware redundancy**
 - Take a distributed system with **4 file servers**, each with a 0.95 chance of being up at any instant
 - The probability of all 4 being down simultaneously is $0.054 = 0.000006$
 - So the probability of at least one being available (i.e., the reliability of the full system) is 0.999994, far better than 0.95
 - If there are 2 servers, then the reliability of the system is $(1 - 0.052) = 0.9975$
 - **Software redundancy**
 - **Process redundancy** with similar considerations
- A design that does not require simultaneous functioning of a substantial number of critical components.

Hardware Redundancy

- Two computers are employed for a single application, one acting as a standby
 - Very costly, but often very effective solution
- Redundancy can be planned at a finer grain
 - Individual servers can be replicated
 - Redundant hardware can be used for non-critical activities when no faults are present
 - Redundant routes in network

Failure Masking by Redundancy


- Triple modular redundancy.




Process Recovery

- Process groups
 - All members of a group receive a message to the group
 - If one process fails, others can take over
 - Can be dynamic; processes can have multiple memberships
 - Flat vs. hierarchical groups

Flat Groups versus Hierarchical Groups



(a)



(b)

- a) Communication in a flat group.
- b) Communication in a simple hierarchical group

Management of Replicated Processes



- Primary copy
 - Primary-backup setup
 - Coordinator is the primary that coordinates all updates
 - If coordinator fails, one backup takes over (usually through an election procedure)
 - Processes are organized hierarchically
- Replicated-writes
 - Active replication and quorum-based protocols
 - Flat group organization
 - No single points of failure

Process resilience

How can **fault tolerance** be achieved in distributed systems?

1. **Key approach** to tolerating a faulty process:

replicate the process and organize these identical process into a group.



Some design issues

- Structure of group: flat/hierarchical
- Need for **managing** groups and group membership
 - **centralized**: group server
 - **distributed**: totally-ordered reliable multicast
- How many replicas? For **K fault** tolerant,
 - Fail-silent faults : $K+1$
 - Byzantine faults : $2K+1$ **majority**

Agreement in faulty system


Introduction

There is a need to have processes **agree on something**.


- **Goal**: all **non-faulty processors reach consensus** on some issue within a finite number of steps
- Two kinds of fault:
 - **communication fault** : two-army problem
 - **processor fault** : Byzantine generals problem

Two-army problem

- Problem:



Two blue armies want to coordinate their attacks on the red army. But they can **only send messenger**



A and B will never reach agreement


Because sender of the last message **doesn't know** if the **last message arrived**.

- Conclusion

agreement between two processes is **not possible** in the case **unreliable communication**

- How does TCP deal with this problem when two computer want to establish a **TCP connection** ?


Three way handshake



Byzantine agreement problem

- Communication is perfect but the **processors are not** (in Byzantine faults) .
- Problem:
 n blue generals want to coordinate their attacks on the red army. But m of them are traitors.

Byzantine agreement problem



Question:

Whether the loyal generals can still reach **agreement**?

Byzantine agreement problem

Generals **exchange troop strengths**, at the end, each **general has a vector** of length n corresponding to **all the armies**.

Let $n=4$, $m=1$


general 1 has 1K troop

general 2 has 2K troop

general 3 is traitor

general 4 has 4K troop

1 Got(1, 2, x, 4) 2 Got(1, 2, y, 4)



3 Got(1, 2, 3, 4) 4 Got(1, 2, z, 4)

- | | |
|------------------|--------------------------|
| 1 is not sure if | 2 sends him true message |
| 1 is not sure if | 3 sends him true message |
| 1 is not sure if | 4 sends him true message |

1 Got
 (1, 2, y, 4)
 (a, b, c, d)
 (1, 2, z, 4)

2 Got
 (1, 2, x, 4)
 (e, f, g, h)
 (1, 2, z, 4)

4 Got
 (1, 2, x, 4)
 (1, 2, y, 4)
 (i, j, k, l)

P1

P2

P3

P4


step 2: (1, 2, x, 4) (1, 2, y, 4) (1, 2, 3, 4) (1, 2, z, 4)

step 3: (1, **2**, y, 4) (**1**, 2, x, 4) (**1**, 2, x, 4) (**1**, 2, x, 4)


(a, b, **c**, d) (e, f, **g**, h) (1, **2**, y, 4) (1, **2**, y, 4)

(1, 2, z, **4**) (1, 2, z, **4**) (1, 2, z, **4**) (i, j, **k**, l)

step 4: (1, 2, u, 4) (1, 2, u, 4) ~~(1, 2, u, 4)~~ (1, 2, u, 4)



Agreement in Faulty Systems (1)



1 Got(1, 2, x, 4)
 2 Got(1, 2, y, 4)
 3 Got(1, 2, 3, 4)
 4 Got(1, 2, z, 4)

(b)

1 Got	2 Got	4 Got
(1, 2, y, 4)	(1, 2, x, 4)	(1, 2, x, 4)
(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)

(c)

- The Byzantine generals problem for 3 loyal generals and 1 traitor.
- a) The generals announce their **troop strengths** (in units of 1 kilosoldiers).
- b) The **vectors** that each general assembles based on (a)
- c) The **vectors** that each general receives in step 3.

- **Algorithm** to reach agreement. They perform the following :

step 1: **every general sends** a message to every other general telling his strength (true or lie)

step 2: each general collects **received** messages to form a vector


step 3: every general **passes his vector** to every other general

step 4: each general **examines the ith element** of each of the **newly received** vectors. If any value has a majority, that value is put into the result vector

- $\textcircled{3}$ is Byzantine faulty processor

A system with m faulty processors, agreement can be achieved only if $2m+1$ processors work properly, for a total of $3m+1$. i.e. $>2n/3$

For example, let $n=3$, $m=1$



P1

P2


P3

step 2: (1, 2, x) (1, 2, y) (1, 2, 3)

step 3: (1, 2, y) (1, 2, x) (1, 2, x)
 (a, b, c) (d, e, f) (1, 2, y)

step 4: (u, u, u) (u, u, u,) ~~(1, 2, u)~~

Agreement in Faulty Systems (2)



(b)

1	Got(1, 2, x)
2	Got(1, 2, y)
3	Got(1, 2, 3)

(c)

1 Got	2 Got
(1, 2, y)	(1, 2, x)
(a, b, c)	(d, e, f)

- The same as in previous slide, except now with 2 loyal generals and one traitor.

Interactive Consistency (IC)


- The question is whether for **given** $m, n > 0$, it is possible to devise an algorithm based on an exchange of messages that will allow each nonfaulty processors to compute **a vector of values** with an element for each of the n processors, such that
 - (1) the **nonfaulty processors** compute exactly the **same vector**;
 - (2) the element of this vector corresponding to a given **nonfaulty processor** is the **private value** of that processor.

1 Got	2 Got	4 Got
(1, 2, y, 4)	(1, 2, x, 4)	(1, 2, x, 4)
(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)
(1, 2, u, 4)	(1, 2, u, 4)	(1, 2, u, 4)

Byzantine Generals Problem

A **commanding general** must send an order to his $n - 1$ **lieutenant** generals such that


- IC1. All **loyal lieutenants** **obey** the same order.
- IC2. If the **commanding** general is **loyal**, then every **loyal lieutenant** obeys the order he sends.



Oral Message Algorithm

Algorithm OM(0):

1. Commander sends his value to every lieutenant
2. Each lieutenant uses the value received or "retreat" if no value received



Algorithm OM(m), $m > 0$:

1. Commander sends his value to every lieutenant
2. For each i , let v_i be the value that lieutenant i receives from the commander or "retreat". Lieutenant i acts as the commander in OM(m-1) to send the value v_i to each of the other $n-2$ other lieutenants
3. For each i , and each $j \neq i$, let v_j be the value lieutenant i received from lieutenant j in step 2. Lieutenant i uses the value *majority* (v_1, \dots, v_{n-1})

