

Object-based Distributed Systems

殷亚凤

Email: yafeng@nju.edu.cn

Homepage: <http://cs.nju.edu.cn/yafeng/>

Room 301, Building of Computer Science and Technology

Review

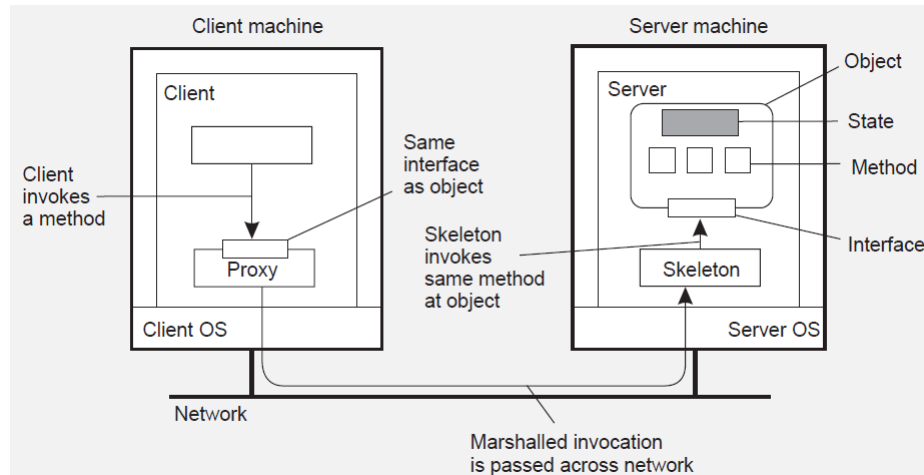
- Introduction
- System Architecture
- Processes
- Communication
- Naming
- Synchronization
- Consistency & Replication
- Fault Tolerance

This Lesson

- Remote distributed objects
- Processes: Object servers
- Remote Method Invocation (RMI)
- Object references
- Consistency and replication

Remote distributed objects

- Data and operations encapsulated in an **object**
- Operations implemented as **methods** grouped into interfaces
- Object offers only its **interface** to clients
- Object server is responsible for **a collection of objects**
- **Client stub** (proxy) implements interface
- **Server skeleton** handles (un)marshaling and object invocation



Remote distributed objects

- Types of objects I
 - **Compile-time objects**: Language-level objects, from which proxy and skeletons are automatically generated.
 - **Runtime objects**: Can be implemented in any language, but require use of an object adapter that makes the implementation appear as an object.
- Types of objects II
 - **Persistent objects**: live independently from a server: if a server exits, the object's state and code remain (passively) on disk.
 - **Transient objects**: live only by virtue of a server: if the server exits, so will the object.

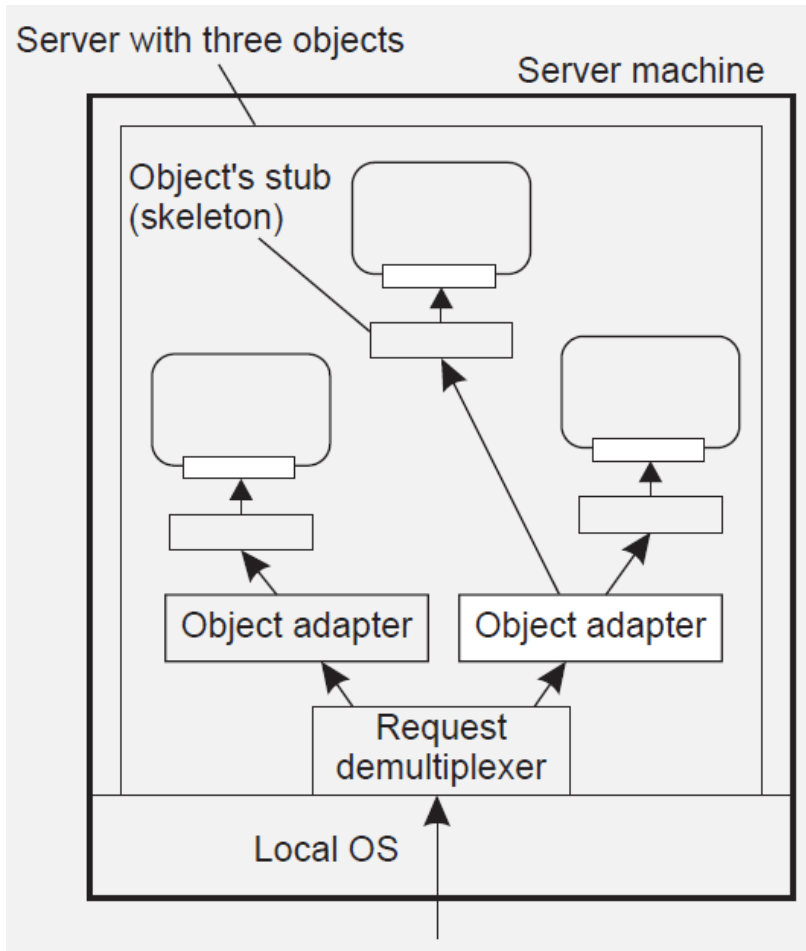
Processes: Object servers

- The actual **implementation of an object**, sometimes containing only method implementations:
 - Collection of C or COBOL functions, that act on structs, records, database tables, etc.
 - Java or C++ classes
- **Server-side stub** for handling network I/O:
 - Unmarshalls incoming **requests**, and calls the appropriate servant code
 - Marshalls results and sends **reply** message
 - Generated from **interface specifications**

Processes: Object servers

- The “manager” of a set of objects:
 - Inspects (as first) **incoming requests**
 - Ensures referenced object is **activated** (requires identification of servant)
 - Responsible for **generating object references**
 - Passes request to **appropriate skeleton**, following specific activation policy

Processes: Object servers



- Object servers determine how their objects are constructed

Example: Ice

```
main(int argc, char* argv[]) {  
    Ice::Communicator ic;  
    Ice::ObjectAdapter adapter;  
    Ice::Object object;  
    ic = Ice::initialize(argc, argv);  
  
    adapter = ic->createObjectAdapterWithEndpoints  
        ( "MyAdapter", "tcp -p 10000");  
    object = new MyObject;  
  
    adapter->add(object, objectID);  
    adapter->activate();  
  
    ic->waitForShutdown();  
}
```

- Activation policies can be changed by modifying the properties attribute of [an adapter](#). Ice aims at simplicity, and achieves this partly by putting policies into the middleware.

Remote Method Invocation (RMI)

- **Assume client stub and server skeleton are in place**
- **Client** invokes method **at stub**
- **Stub** marshals request and **sends it to server**
- **Server** ensures **referenced object is active**:
 - Create separate process to **hold object**
 - Load the object **into server process**
 - ...
- Request is **unmarshaled by object's skeleton**, and referenced method is invoked
- If request **contained an object reference**, invocation is applied recursively (i.e., server acts as client)
- Result is **marshaled** and passed back **to client**
- **Client stub unmarshals reply** and passes result **to client application**

RMI: Parameter passing

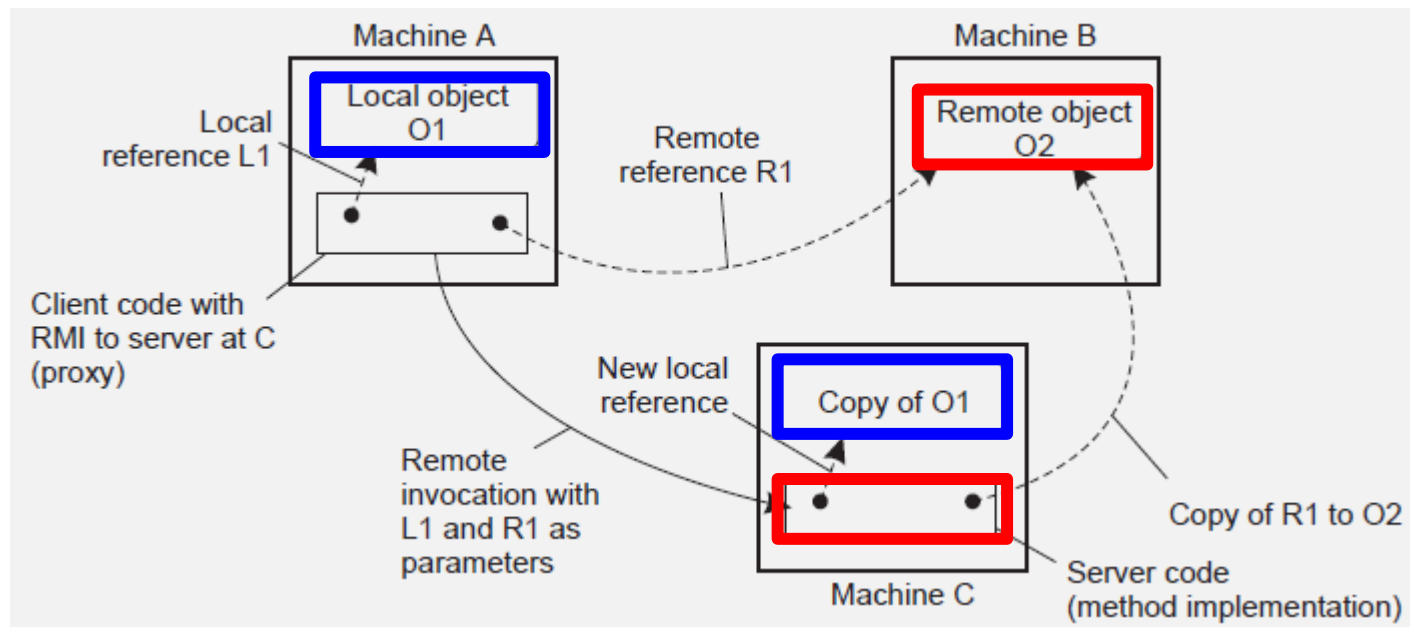
- Much easier than in the case of RPC:
 - Server can simply **bind to referenced object**, and invoke methods
 - **Unbind** when referenced object is no longer needed

RMI: Parameter passing

- A client may also pass a complete object as parameter value:
 - An object has to be marshaled:
 - Marshall its state
 - Marshall its methods, or give a reference to where an implementation can be found
 - Server unmarshals object. Note that we have now created a copy of the original object.
 - Object-by-value passing tends to introduce nasty problems

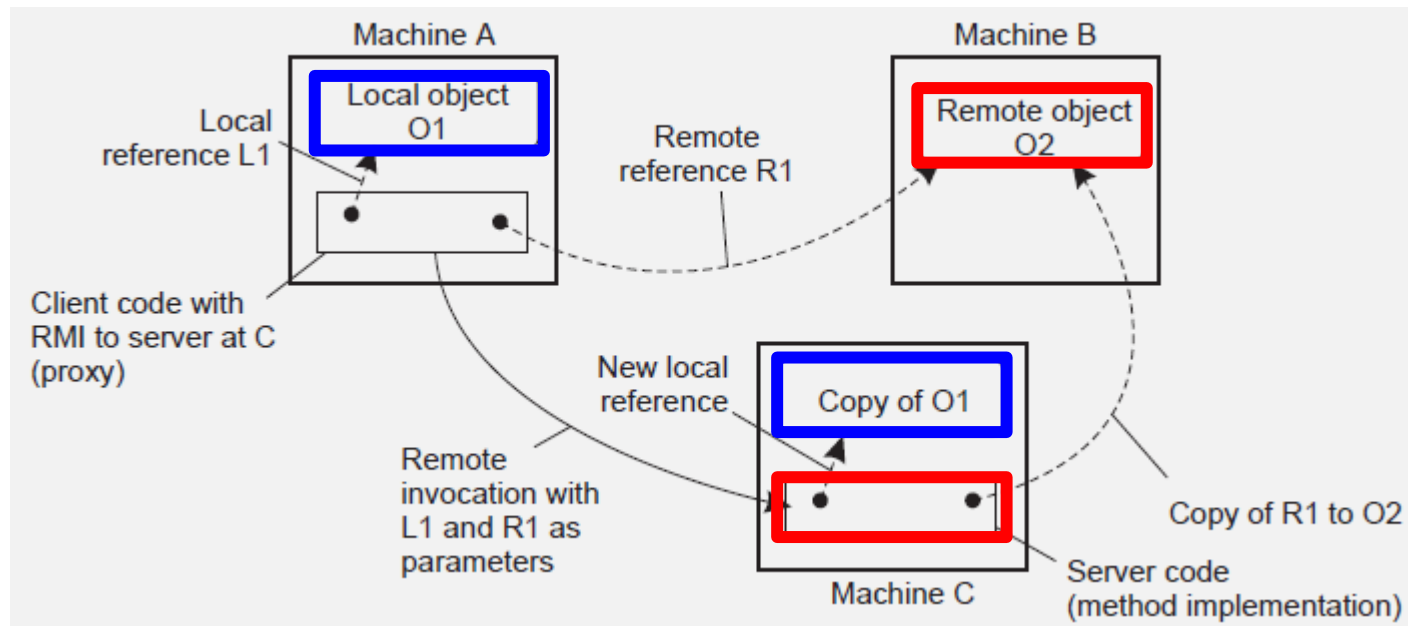
RMI: Parameter passing

- Systemwide object reference generally contains **server address, port** to which adapter listens, and **local object ID**.
Extra: Information on protocol between client and server (TCP, UDP, SOAP, etc.)

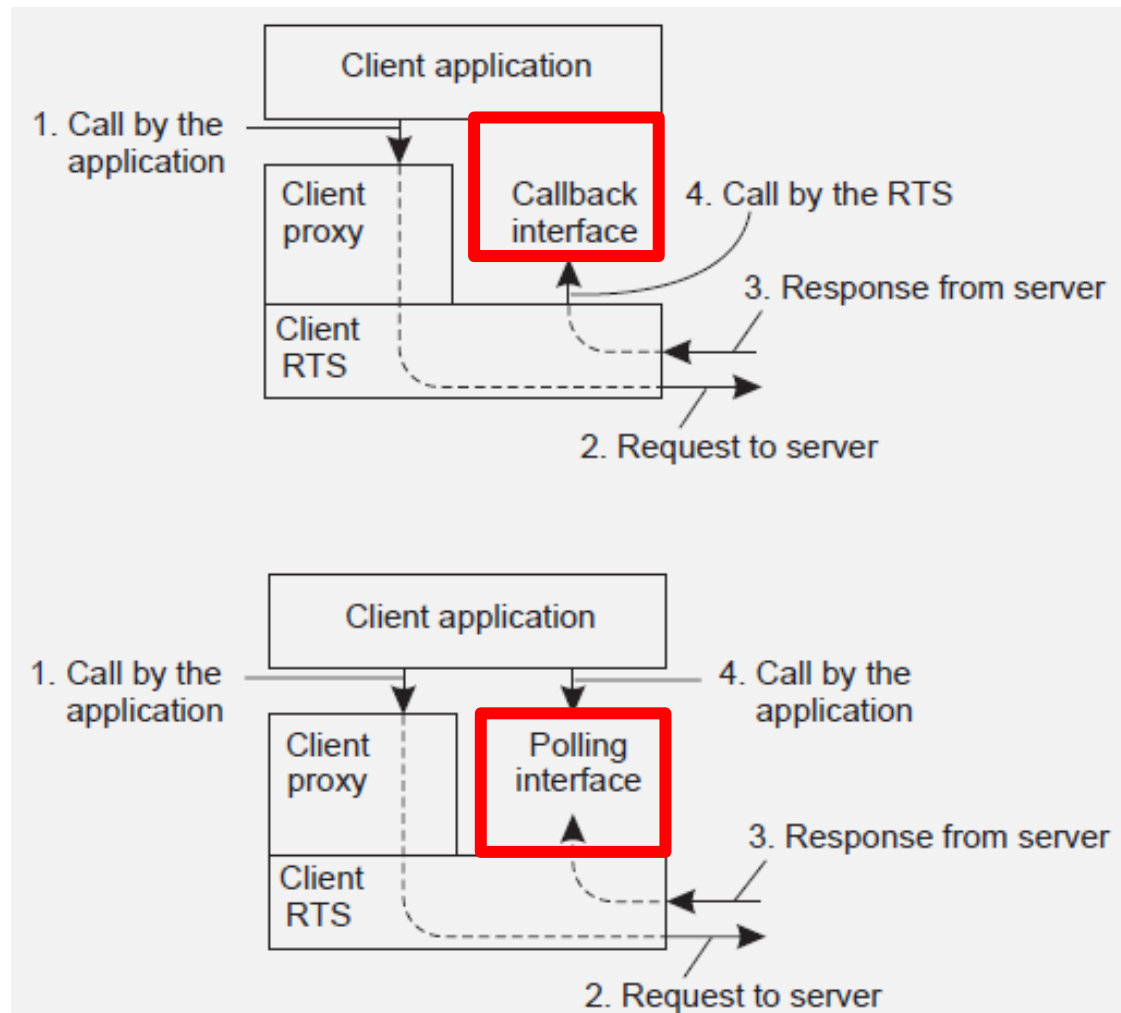


RMI: Parameter passing

- Question: What's an alternative implementation for a remote-object reference?

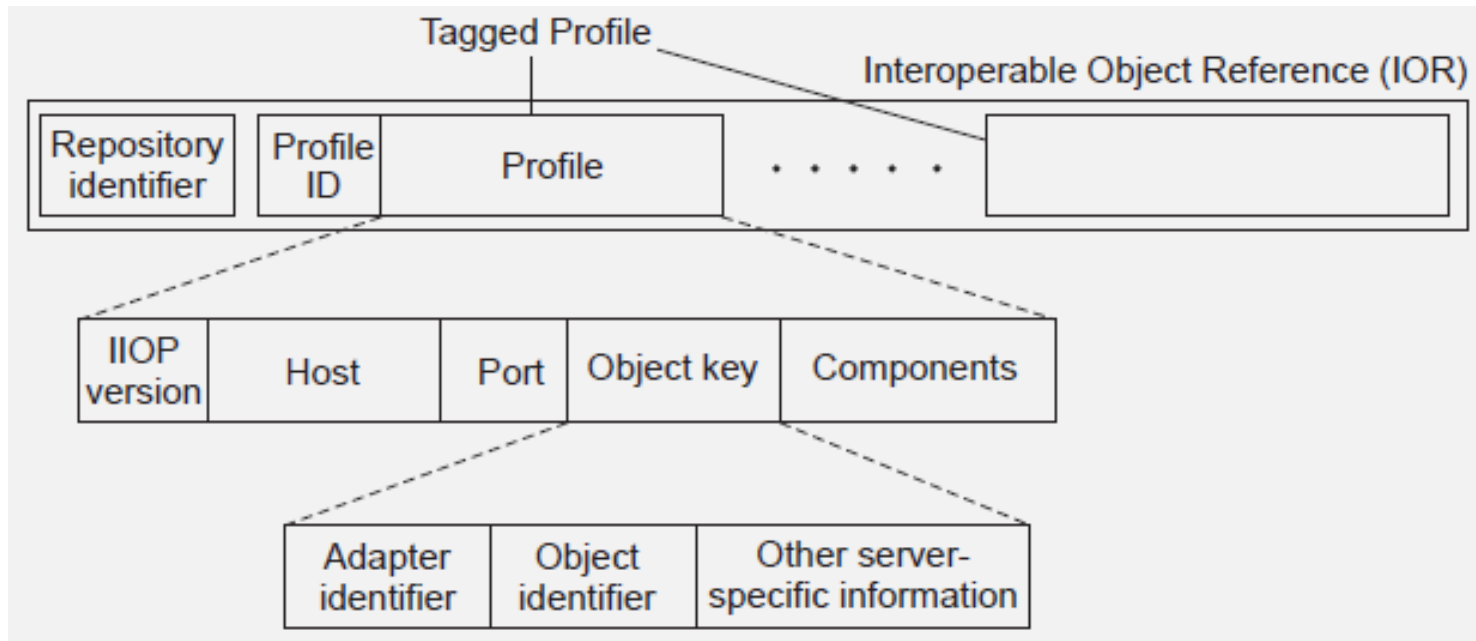


Object-based messaging



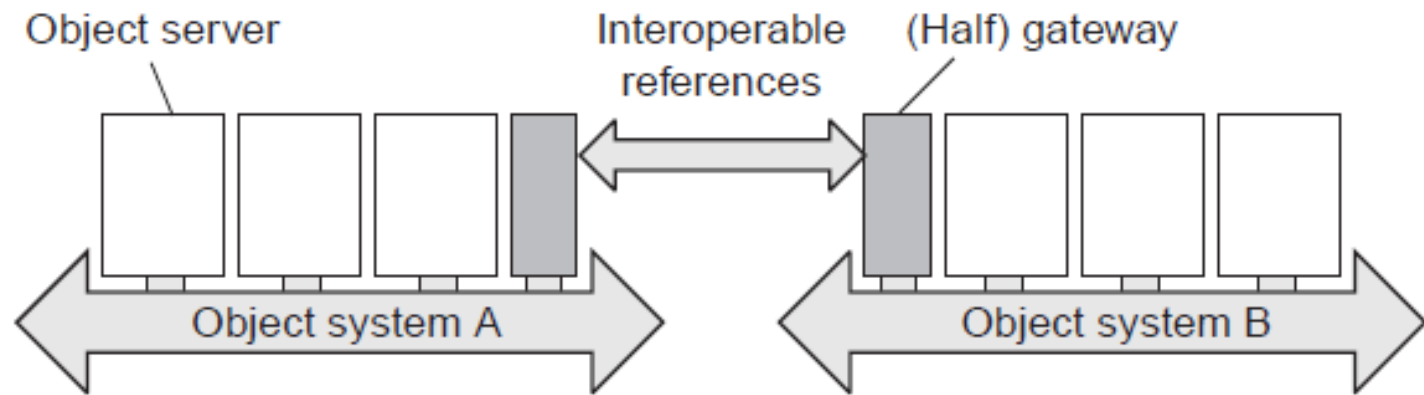
Object references

- In order to **invoke remote objects**, we need a means to uniquely refer to them. **Example:** CORBA object references.



Object references

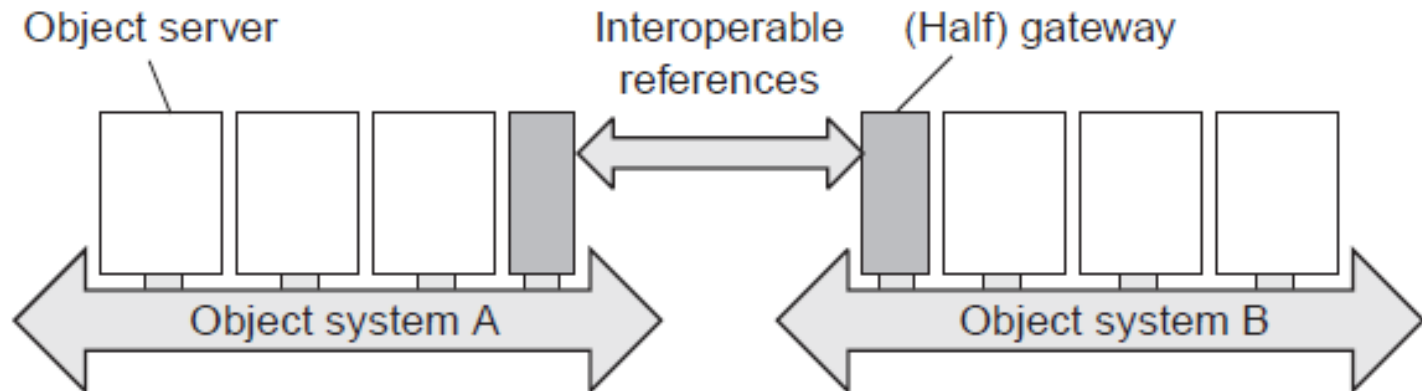
- It is not important how object references are implemented per object-based system, as long as there is **a standard to exchange them** between systems.



- Object references passed from one RTS to another are transformed **by the bridge** through which they pass (different transformation schemes can be implemented)

Object references

- Passing an object reference *refA* from RTS A to RTS B circumventing the A-to-B bridge may be useless if RTS B doesn't understand *refA*



Globe object references: location independent

- **Stack of addresses** representing the protocol to speak:

Field	Description
Protocol ID	Constant representing a (known) protocol
Protocol addr.	Protocol-specific address
Impl. handle	Reference to a file in a repository

- Contains all that is needed to talk in **a proprietary way** to an object:

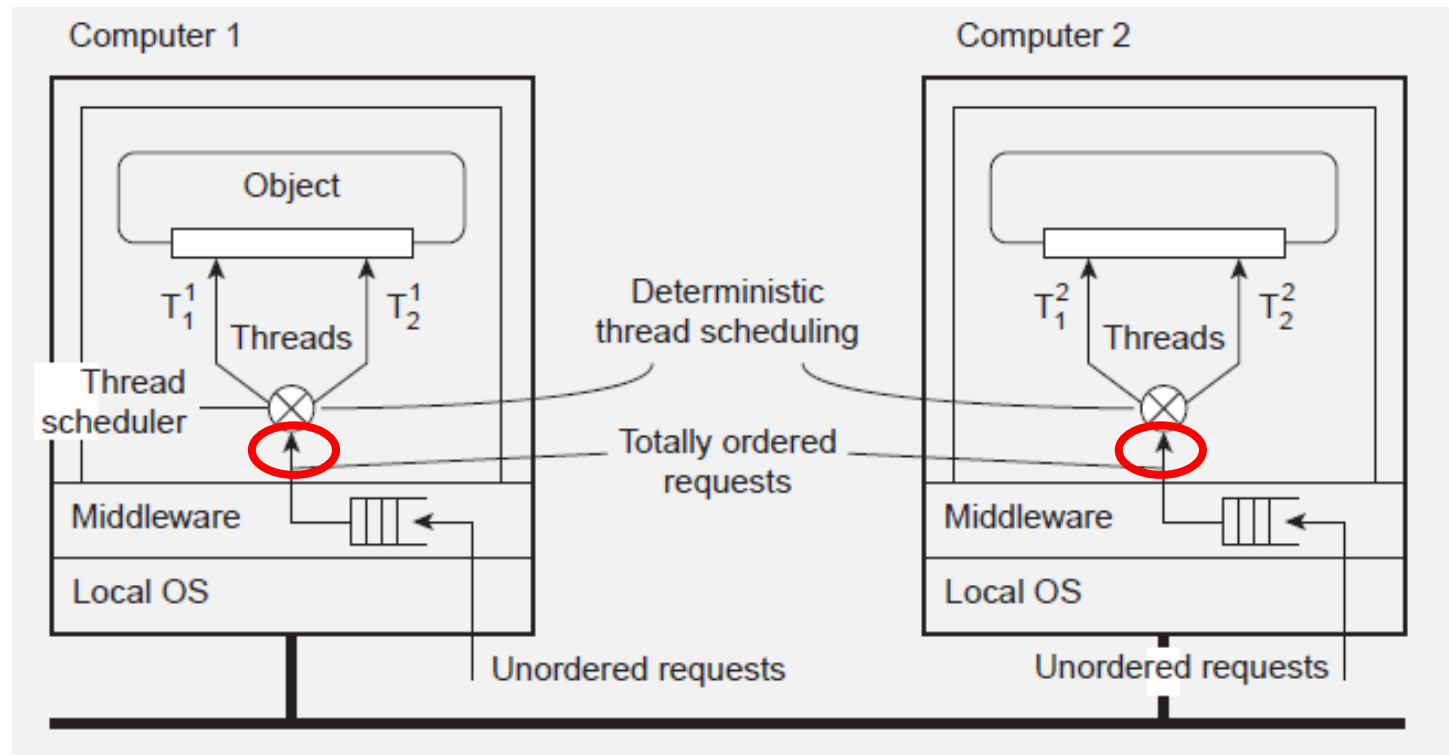
Field	Description
Impl. handle	Reference to a file in a repository
Initialization string	Used to initialize an implementation

Consistency and replication

- Objects form a natural means for realizing **entry consistency**:
 - Data are grouped into units, and protected by a synchronization variable (i.e., lock)
 - Synchronization variables adhere to sequential consistency (i.e., values are set atomically)
 - Operations of grouped data can be nicely grouped: object
- What happens when objects are **replicated**? One way or the other we need to ensure that operations on replicated objects are properly ordered.

Replicated objects

- We need to make sure that **requests are ordered correctly** at the servers and that threads are deterministically scheduled

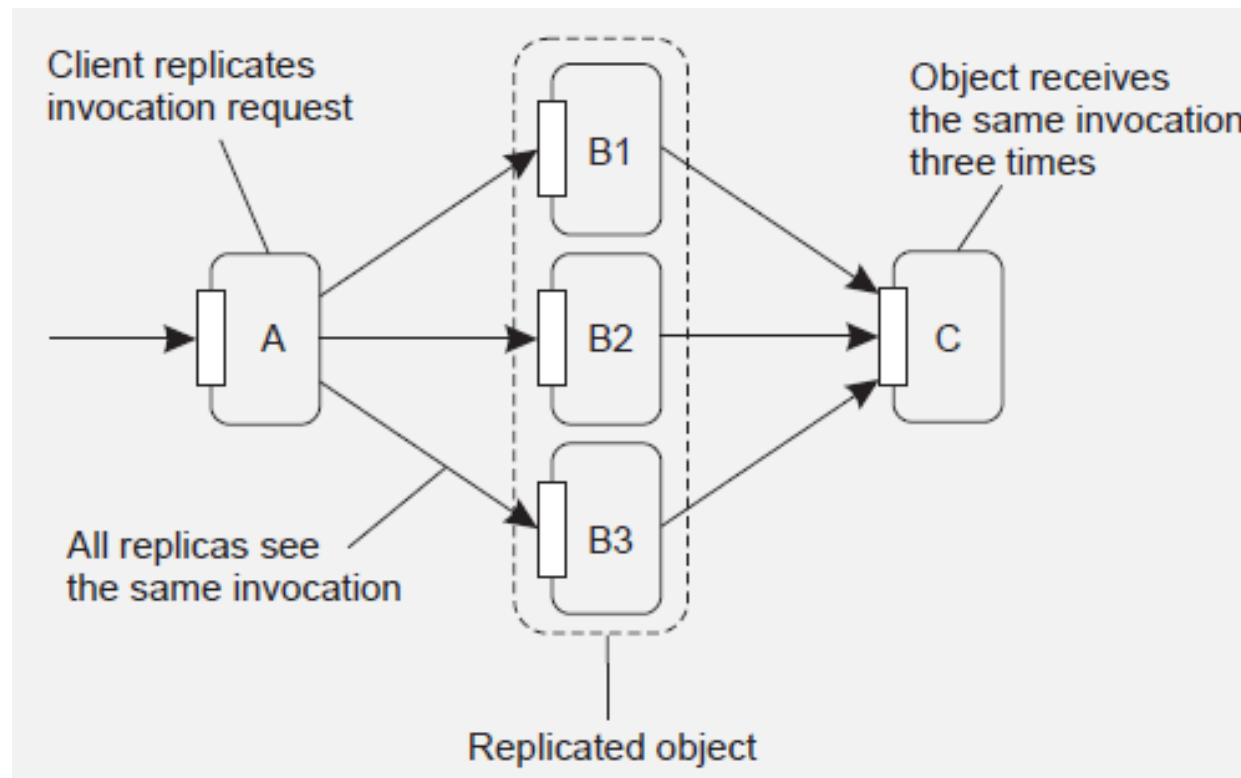


Replicated objects

- We are dealing with nasty issues here. Simplicity may dictate completely serialized (i.e., single-threaded) executions at the server.

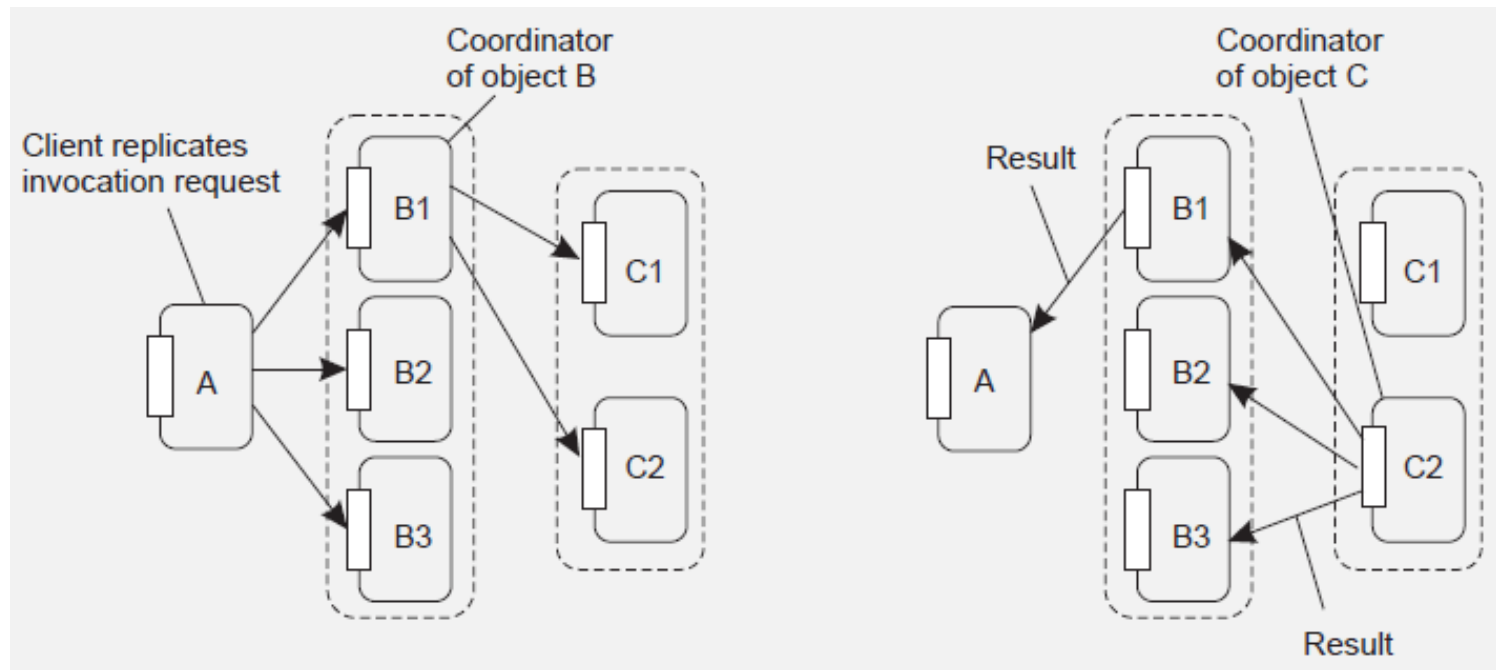
Replicated invocations

- Updates are forwarded to **multiple replicas**, where they are carried out. There are some problems to deal with in the face of **replicated invocations**



Replicated invocations

- Assign a **coordinator on each side** (client and server), which ensures that only one invocation, and one reply is sent



Assignment 2 & 3

- <https://www.cs.princeton.edu/courses/archive/fall16/cos418/a2.html> (Go Language)
- Alternative: determine what you will do by yourself
- Presentation (5-8 minutes): June 3, 2020 / June 10, 2020
- Code and Report: before June 30, 2020
- Submission: distrisys@126.com