

Theoretical Models of MapReduce-type Computation

张晓达

Paper Presentations

- Howard Karloff, Siddharth Suri, Sergei Vassilvitskii. **A Model of Computation for MapReduce. *SODA '10*.**
- Tim Roughgarden, Sergei Vassilvitskii, and Joshua R. Wang. **Shuffles and Circuits: (On Lower Bounds for Modern Parallel Computation). *SPAA '16*. [Best Paper Award]**
- Sungjin Im, Benjamin Moseley, and Xiaorui Sun. **Efficient Massively Parallel Methods for Dynamic Programming. *STOC '17*.**

Outline

- Background
 - Traditional PRAM and **NC**
 - MapReduce
- The fundamental problems introduced by MapReduce
- The model and a particular problem [KSV *SODA* '10]
- The model and its *lower bound* [RVW *SPAA* '16]
- The model and a specific *algorithm design technique* [IMS *STOC* '17]

Background

- The synchronous parallel random access machine (**PRAM**) model
 - Global memory (M locations), P processors (each has a constant number of local registers)
 - Computation proceeds in a series of *synchronous* parallel steps, each consisting:
 - each processor chooses a mem. location and reads;
 - each does some computation;
 - each chooses a mem. location and writes;
 - By *synchrony*, every processor finishes executing step i before any processor begins executing step $i+1$
- Whether or not allowing concurrent read/write
 - EREW, CREW, CRCW

Background

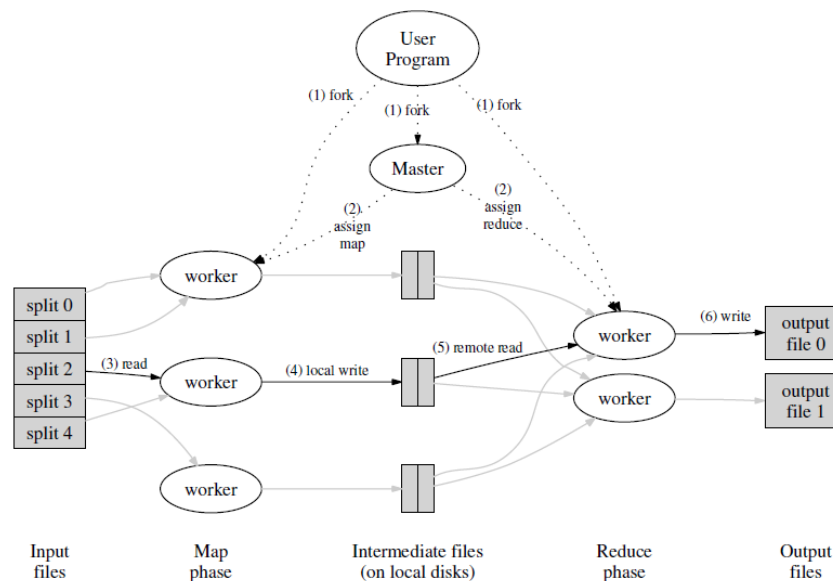
- What problems can be solved in *polylog* time on a PRAM with a *polynomial* number of processors?
 - Polylog time serves as *gold-standard* for parallel running time;
 - Polynomial number of processors provides a necessary condition for *efficiency*.
- The class **NC** (Nick's class) is defined as the set of such problems

► **Definition 12.1:** The class *NC* consists of languages *L* that have a PRAM algorithm *A* such that for any $x \in \Sigma^*$

- $x \in L \Rightarrow A(x)$ accepts;
- $x \notin L \Rightarrow A(x)$ rejects;
- the number of processors used by *A* on *x* is polynomial in $|x|$;
- the number of steps used by *A* on *x* is polylogarithmic in $|x|$.

Background

- MapReduce
 - Map phase, (shuffle), Reduce phase
 - Tasks in Map phase compute over disjoint input splits and output intermediate data *locally*
 - Tasks in Reduce phase take as input the intermediate data, and contribute the parts of the final output



The fundamental problems

- What is *efficiently* (e.g. in polylog rounds) computable in the MapReduce paradigm?
- What is the fundamental *limitations* on how efficiently algorithms implemented on MapReduce?
- Is there a *meta-algorithm* that takes as input a sequential algorithm and outputs an efficient distributed (MapReduce) algorithm?

The MRC

The input to a **MRC** problem is a finite sequence of pairs $\langle k_j, v_j \rangle$, for $j = 1, 2, 3 \dots$. The length of the input is $\bar{n} = \sum_j (|k_j| + |v_j|)$.

randomized
version

DEFINITION Fix an $\epsilon > 0$. An algorithm in \mathcal{MRC}^i consists of a sequence $\langle \mu_1, \rho_1, \mu_2, \rho_2, \dots, \mu_R, \rho_R \rangle$ of operations which outputs the correct answer with probability at least $3/4$ where:

number of machines
is sublinear $O(n^{1-\epsilon})$

sublinear
space in each
machine

- Each μ_r is a randomized mapper implemented by a RAM with $O(\log n)$ -length words, that uses $O(n^{1-\epsilon})$ space and time polynomial in n .
- Each ρ_r is a randomized reducer implemented by a RAM with $O(\log n)$ -length words, that uses $O(n^{1-\epsilon})$ space and time polynomial in n .
- The total space $\sum_{\langle k;v \rangle \in U_r} (|k| + |v|)$ used by $\langle \text{key}; \text{value} \rangle$ pairs output by μ_r is $O(n^{2-2\epsilon})$.
- The number of rounds $R = O(\log^i n)$.

允许mapper阶段的
输出有duplication,
但不是任意
duplication

Restriction is only applied to
mapper, not reducer.

- (deterministic **MRC**) **DMRC**'s relationship with **P/NC**

- **DMRC** is in **P**

THEOREM If $P \neq NC$ then $DMRC \not\subseteq NC$.

PROOF: From a **P-complete** language (Circuit Value), construct another language which is solvable in **DMRC**. Since **P-complete** language is not in **NC**, **DMRC** is not in **NC**.

- **DMRC** ?= **P**

DMRC实际定义的language是solvable by Turing machine simultaneously in space $O(n^2 \log n)$ and polynomial time。
若证明**P** \neq **MRC**，需找到一个language在**P - MRC**中，这样的language目前还不知道是否存在。

MRC-parallelizable function

- MRC-parallelizable function as a building block:

DEFINITION *Let S be a set. Call a function f on S MRC-parallelizable if there are functions g and h so that:*

1. *For any partition $\mathcal{T} = \{T_1, T_2, \dots, T_k\}$ of S , where $\cup_i T_i = S$ and $T_i \cap T_j = \emptyset$ for $i \neq j$ (of course), f can be expressed as: $f(S) =$*

$h(g(T_1), g(T_2), \dots, g(T_k))$.

2. *g and h can be expressed in $O(\log n)$ bits.*
3. *g and h can be computed in time polynomial in $|S|$ and every output of g can be expressed in $O(\log n)$ bits.*

applying g ,
and then h

- If one wants to compute f over S , one could do so by partitioning S arbitrarily, applying g to each part of the partition, and then applying h to the results.

The power of **MRC**-parallelizable function


LEMMA *Consider a universe \mathcal{U} of size n and a collection $\mathcal{S} = \{S_1, \dots, S_k\}$ of subsets of \mathcal{U} , where $S_i \subseteq \mathcal{U}$, $\sum_{i=1}^k |S_i| \leq n^{2-2\epsilon}$, and $k \leq n^{2-3\epsilon}$. Let $\mathcal{F} = \{f_1, \dots, f_k\}$ be a collection of **MRC**-parallelizable functions. Then the output $f_1(S_1), \dots, f_k(S_k)$ can be computed using $O(n^{1-\epsilon})$ reducers each with $O(n^{1-\epsilon})$ space.*

- 构造性证明: (两轮MapReduce)需要split 每个 S_i , 让g先对每个partition进行计算, 而后再用h再进行一次计算可得最后结果。
- The power of this lemma: **focus on** the structure of the problem and the input; the lemma will handle **how to distribute** the input across the reducers

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

Application of the Function lemma

- To solve the USTCON problem:

1. Begin with every node $v \in V$ being active with label $\ell(v) = v$.
2. For $i = 1, 2, 3, \dots, O(\log N)$ do:
 - (a) Call each active node a leader with probability $1/2$.
 - (b) For every active non-leader node w , find the smallest (according to π) node $w^* \in \Gamma'(w)$.
 - (c) If w^* is not empty, mark w passive and relabel each node with label w by w^* .
3. Output *true* if s and t have the same labels, false otherwise.




- This Monte Carlo algorithm is correct
- It needs $O(\log N)$ rounds to terminate, w.h.p.
- The (a)(b)(c) step *satisfies* the conditions of the **Function Lemma**

The s-shuffle model and lower bound

- **s-shuffle model**

- few rounds s-shuffle computation and its connection to low-degree polynomials
 - translate a lower bound on the *polynomial degree* of a Boolean function to a lower bound of *s-shuffle computation*
- prove lower bounds on the polynomial degree (Boolean)



The s-shuffle model

Capturing core properties of massively parallel computation

- Computation proceeds in **synchronous** rounds
- In each round, each machine performs an arbitrary computation on its input, and send arbitrary information to arbitrary machines in the next round (**shuffle**)
- The only constraint is that each machine receives **at most s bits** each round

Definition (**s -SHUFFLE Computation**) An R -round s -SHUFFLE computation with inputs x_1, \dots, x_n and outputs y_1, \dots, y_k has the following ingredients:


1. A set V of *machines*, which includes one machine for each input bit x_i and each output bit y_i .
2. An assignment of a *round* $r(v)$ to each machine $v \in V$. Machines corresponding to input bits have round 0. Machines corresponding to output bits have round $R + 1$. All other machines have a round in $\{1, 2, \dots, R\}$.
3. For each pair (u, v) of machines with $r(u) < r(v)$, a function α_{uv} from $\{0, 1, \perp\}^s$ to $\{0, 1, \perp\}^s$.

The s-shuffle model

Definition (Result of an s -SHUFFLE Computation) The *result* of an s -SHUFFLE computation assigns a value $g(v) \in \{0, 1, \perp\}^s$ to every machine $v \in V$, and is defined inductively as follows.

1. For a round-0 machine v , corresponding to an input bit x_i , the value $g(v)$ is the s -tuple $(x_i, \perp, \perp, \dots, \perp)$.
2. Given the value $g(u)$ assigned to every machine u with $r(u) < q$, the value assigned to a machine v with $r(v) = q$ is the \perp -sum, over all machines u with $r(u) < r(v)$, of the message $\alpha_{uv}(g(u))$ sent to v by u :

$$g(v) := \odot_{u : r(u) < r(v)} \alpha_{uv}(g(u)).$$



The main theorem


Theorem 3.1 *Suppose that an s -SHUFFLE computation computes the function $f : \{0, 1\}^n \rightarrow \{0, 1\}^k$ in r rounds. Then there are k polynomials $\{p_i(x_1, \dots, x_n)\}_{i=1}^k$ of degree at most s^r such that $p_i(\mathbf{x}) = f(\mathbf{x})_i$ for all $i \in \{1, 2, \dots, k\}$ and $\mathbf{x} \in \{0, 1\}^n$.*

PROOF: Induction on the number of rounds.

We claim that for every non-output machine $v \in V$, value $\mathbf{z} \in \{0, 1, \perp\}^s$, there is a polynomial $p_{v,\mathbf{z}}(x_1, \dots, x_n)$ that evaluates to 1 on point \mathbf{x} for which the computation's assigned value $g(v)$ to v is \mathbf{z} and to 0 on all other points $\mathbf{x} \in \{0, 1, \perp\}^n$. $p_{v,\mathbf{z}}$ has degree at most $s^{r(v)}$.

Base case: for a machine v in round 0 ($r(v) = 0$). Each such machine corresponds to an input bit x_i and its value $g(v)$ is $(x_i, \perp, \perp, \dots, \perp)$. The (degree-1) polynomial for $\mathbf{z} = (0, \perp, \dots, \perp)$ and $\mathbf{z} = (1, \perp, \dots, \perp)$ are $p_{v,\mathbf{z}}(x_1, \dots, x_n) = 1 - x_i$ and $p_{v,\mathbf{z}}(x_1, \dots, x_n) = x_i$, respectively. All other values of \mathbf{z} are impossible for such a machine, so they have polynomials $p_{v,\mathbf{z}}(x_1, \dots, x_n) = 0$.

The main theorem



Theorem 3.1 Suppose that an s -SHUFFLE computation computes the function $f : \{0, 1\}^n \rightarrow \{0, 1\}^k$ in r rounds. Then there are k polynomials $\{p_i(x_1, \dots, x_n)\}_{i=1}^k$ of degree at most s^r such that $p_i(\mathbf{x}) = f(\mathbf{x})_i$ for all $i \in \{1, 2, \dots, k\}$ and $\mathbf{x} \in \{0, 1\}^n$.

Consider a machine v (neither an input bit nor an output bit).

Fix some potential value $\mathbf{z} \in \{0, 1, \perp\}^s$ of $g(v)$, and focus first on the i th coordinate of $g(v)$. Consider some machine u of previous round. It sends z_i on port i to machine v when $\alpha_{\{uv\}}(g(u))$ has an i th entry of z_i .

By inductive hypothesis, for particular value $g(u)$, there is a polynomial of degree at most $s^{r(v)-1}$ that indicates the inputs \mathbf{x} for which u receives this value.


To obtain a polynomial that represents the inputs \mathbf{x} for which $g(v) = \mathbf{z}$, just take the product, $g(v)_1 = z_1, g(v)_2 = z_2, \dots, g(v)_s = z_s$. This produces a polynomial of degree $s^{r(v)}$.

The main theorem

Theorem Suppose that an s -SHUFFLE computation computes the function $f : \{0, 1\}^n \rightarrow \{0, 1\}^k$ in r rounds. Then there are k polynomials $\{p_i(x_1, \dots, x_n)\}_{i=1}^k$ of degree at most s^r such that $p_i(\mathbf{x}) = f(\mathbf{x})_i$ for all $i \in \{1, 2, \dots, k\}$ and $\mathbf{x} \in \{0, 1\}^n$.


逆否命题

Corollary If some output bit of the function $f : \{0, 1\}^n \rightarrow \{0, 1\}^k$ cannot be represented by a polynomial with degree less than d , then every s -SHUFFLE computation that computes f uses at least $\lceil \log_s d \rceil$ rounds.



The lower bound

Theorem *The degree of the Boolean function for undirected ST-CONNECTIVITY on n vertices is $\binom{n}{2}$.*



The model and an algorithm design technique (DP)

- The model

- Let n be the input size, m be the number of machines.
- *local* memory on each machine should be $\tilde{\Theta}(n/m)$, and the algorithm should allow for $n^\epsilon \leq m \leq n^{1-\epsilon}$, for some constant $\epsilon > 0$

Tilde hides polylogarithmic factors

- Are there *principled guidelines* for converting a sequential DP into an efficient (e.g. $O(1)$ rounds) distributed DP?

Key Properties

(1) **Monotonicity**: A sub-problem has no greater (smaller, resp.) optimum if the objective is to be maximized (minimized, resp.). (2)

Decomposibility: The input can be decomposed into a two-level laminar family of partial input, where an upper level partial input is called a group and a lower level partial input is called a block, such that:

- A nearly optimal solution (for the whole input) can be constructed by concatenating solutions for groups.
- A nearly optimal solution for each group can be constructed from $O(1)$ blocks.

Weighted Interval Selection (WIS)

- Definition
 - **Input:** a collection of intervals $\{I_i = (a_i, b_i)\}_{i \in [n]}$, with their respective weights $\{w_i\}_{i \in [n]}$
 - **Goal:** choose a subset of disjoint intervals with the maximum weight.
- Standard sequential DP
 - Assume that intervals are ordered in increasing order of their start points
 - Let **OPT**(S) denote the optimal solution to the instance consisting of a subset S of intervals, (or the optima value). Let **A**(i) denote **OPT**($\{I_i, I_{i+1}, \dots, I_n\}$). Goal is to compute **A**(1) using the following recursion:

Choose the i th interval or not?

No

Yes

$$A(i) = \max\{A(i+1), w_i + A(j)\}$$

where $j = \arg \min_{j'} (b_i < a_{j'})$ when $i \leq n$; and $A(n+1) = 0$.

Preliminaries

- Goal

THEOREM 4.1. For any $\epsilon, \delta > 0$, there exists a $(1 - \epsilon)$ -approximation for the Weighted Interval Selection problem running in $O(\frac{1}{\delta}(\log \frac{1}{\epsilon} + \log \frac{1}{\delta}))$ rounds when each machine has memory $\tilde{O}(n^\delta)$.

- Preprocessing

- 因为最后的结果只取决于interval开始和结束的相对位置，所以可以调整intervals, 使其starting and ending time are all distinct and not integers.
- 将输入均分到m machines上. The first n/m intervals go to machine 1, the next go to machine 2, and so on. Assume that machine k can only see intervals starting at time between k and $k+1$.

The beginning lemma

- Definition
 - block: a subset of disjoint intervals
 - pairwise independent blocks: span disjoint sets of machines
 - L -block: a block contains at most L crossing intervals

LEMMA 4.2. *For any even integer $L \geq 2$, there exists a $(1 - 2/L)$ -approximate solution consisting of (pair-wise) independent L -blocks.*

PROOF: Consider a fixed optimal solution. Partition the crossing intervals into **groups** so that each group contains L consecutive crossing intervals; (the last may contain less).

We **remove the lightest interval** from every group except the last. Clearly, we lose at most $1/L$ times the optimum. It is easy to see that the resulting solution consists of $2L$ -blocks that are pairwise independent.

Overview of the algorithm

- 如果a nearly optimal solution consisting only of 0-blocks 就好了
 - 那么只需要each machine compute the **best ‘local’** solution from the local intervals on the machine, and simply **aggregate** the weights of the local solutions from all machines.
 - The aggregation can be done in $O(1/\delta)$, machine has $\tilde{O}(n^\delta)$ memory.
 - **However**, the optimal solution may have lots of heavy crossing intervals, which can not be ignored.
- By the above lemma
 - We can construct a $(1 - \epsilon)$ -approximate solution from $2/\epsilon$ -blocks that are pairwise independent.

*We compute such blocks approximately and memory-efficiently in $O(1)$ rounds. 一旦我们有了这些blocks, we view each block as an **consolidated interval**. 这样, the new instance whose size is almost linear in the number of machines. 至多只需递归 $O(1/\delta)$ 次*

Formally state the idea

Definition 4.3. For each interval I_i and a target weight μ , let $\text{OPT}(i, \mu, L)$ be the smallest interval index $j \geq i$ such that there is a L -block of weight at least μ that is a subset of $\{I_i, I_{i+1}, I_{i+2}, \dots, I_n\}$ and is disjoint from $\{I_j, I_{j+1}, \dots, I_n\}$. Let $W := \{0, \frac{1}{n^2}, \frac{(1+\eta)}{n^2}, \frac{(1+\eta)^2}{n^2}, \dots, n\}$ for η to be determined. A family $\mathcal{F} = \{D(i, \mu, L)\}_{i \in [n], \mu \in W}$ of blocks is said to be a $1 - \gamma$ -approximate compact family of L -blocks if $D(i, \mu, L)$ has weight at least $(1 - \gamma)\mu$ and $D(i, \mu, L) \leq \text{OPT}(i, \mu, L)$. Here, $D(i, \mu, L)$ is analogously defined as $\text{OPT}(i, \mu, L)$.

- Idea

- there is a $(1-2/L)$ -approximate solution consisting only of L -blocks. 对于 each of those blocks, D^* , the **compact family** contains 几乎和 D^* 一样好的 D , 其中 D 的 **weight is within $(1-\gamma)$ of D^* 's weight and ends no later than D^*** . 因此, 如果我们将 nearly optimal solution 中的 block 换成相应的 the compact family 中的 block。这样得到的 solution 是可行解且近似度是 $(1-2/L)(1-\gamma)$.

Formally state the idea

LEMMA 4.4. *If we can construct a $(1 - O(\epsilon\delta))$ -approximate compact family of L -blocks in $O(\log \frac{1}{\epsilon} + \log \frac{1}{\delta})$ rounds where $L = \frac{2}{\epsilon\delta}$ and $\eta = \frac{\epsilon\delta}{10(\log(1/\epsilon) + \log(1/\delta))}$, then we can obtain a $(1 - O(\epsilon))$ -approximate solution for the IS problem in $O(\frac{1}{\delta} \cdot (\log \frac{1}{\epsilon} + \log \frac{1}{\delta}))$ rounds.*

- It now remains to show how to construct the Desired Compact Family using DP.

find a desired compact family, $\mathcal{F}(\ell) := \{D(i, \mu, 2^\ell - 1)\}_{i \in [n], \mu \in W}$, as stated in Lemma 4.4. Towards this end, we find such a family for $\ell = 0, 1, 2, \dots, \ell_1 = \lceil \log 2/(\epsilon\delta) \rceil$ in this order until we have $2^{\ell} - 1 \geq L = 2/(\epsilon\delta)$. Since we will use an induction on the value of ℓ , we refine and strengthen the property that $\mathcal{F}(\ell)$ satisfies: our goal is to find $\mathcal{F}(\ell)$ that is $(1 - \eta)^\ell$ -approximate in increasing order of $\ell = 0, 1, 2, \dots, \ell_1$ in $O(\ell)$ rounds. Since $L = \frac{2}{\epsilon\delta}$, we only need $O(\log 1/\epsilon + \log 1/\delta)$ rounds, and we have $(1 - \eta)^{\ell_1} = 1 - O(\epsilon\delta)$ since $\eta = \frac{\epsilon\delta}{10(\log(1/\epsilon) + \log(1/\delta))}$.

Our remaining goal is to find $\mathcal{F}(\ell)$ that is $(1 - \eta)^\ell$ -approximate in increasing order of $\ell = 0, 1, 2, 3, \dots, \ell_1$ in $O(\ell)$ rounds.

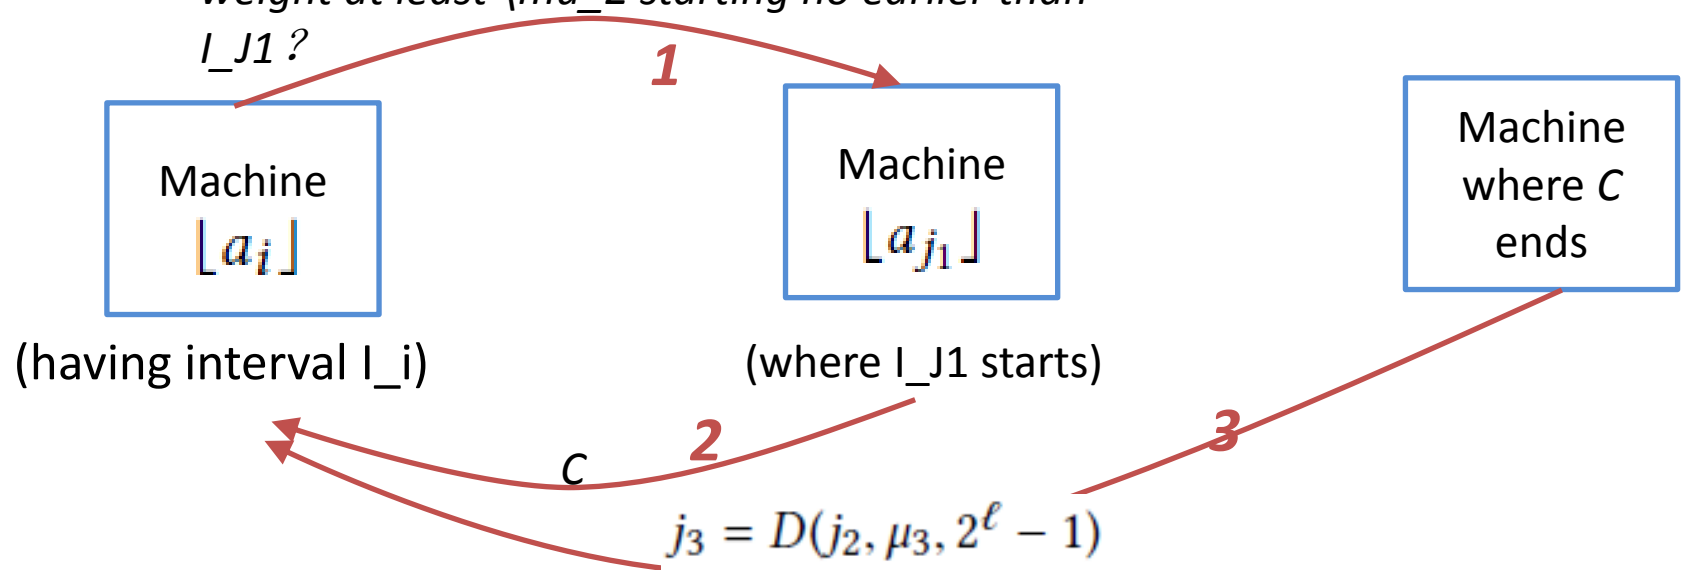
find $F(l+1)$ from $F(l)$ in $O(1)$ rounds.

Show how to compute $D(i, \mu, 2^{\ell+1} - 1)$

Consider any triplet $\mu_1, \mu_2, \mu_3 \in W$ such that

$$(1 - \eta)\mu \leq \mu_1 + \mu_2 + \mu_3 \leq \mu \quad j_1 = D(i, \mu_1, 2^\ell - 1)$$

Which is the earliest ending crossing interval of weight at least μ_2 starting no earlier than I_{j_1} ?

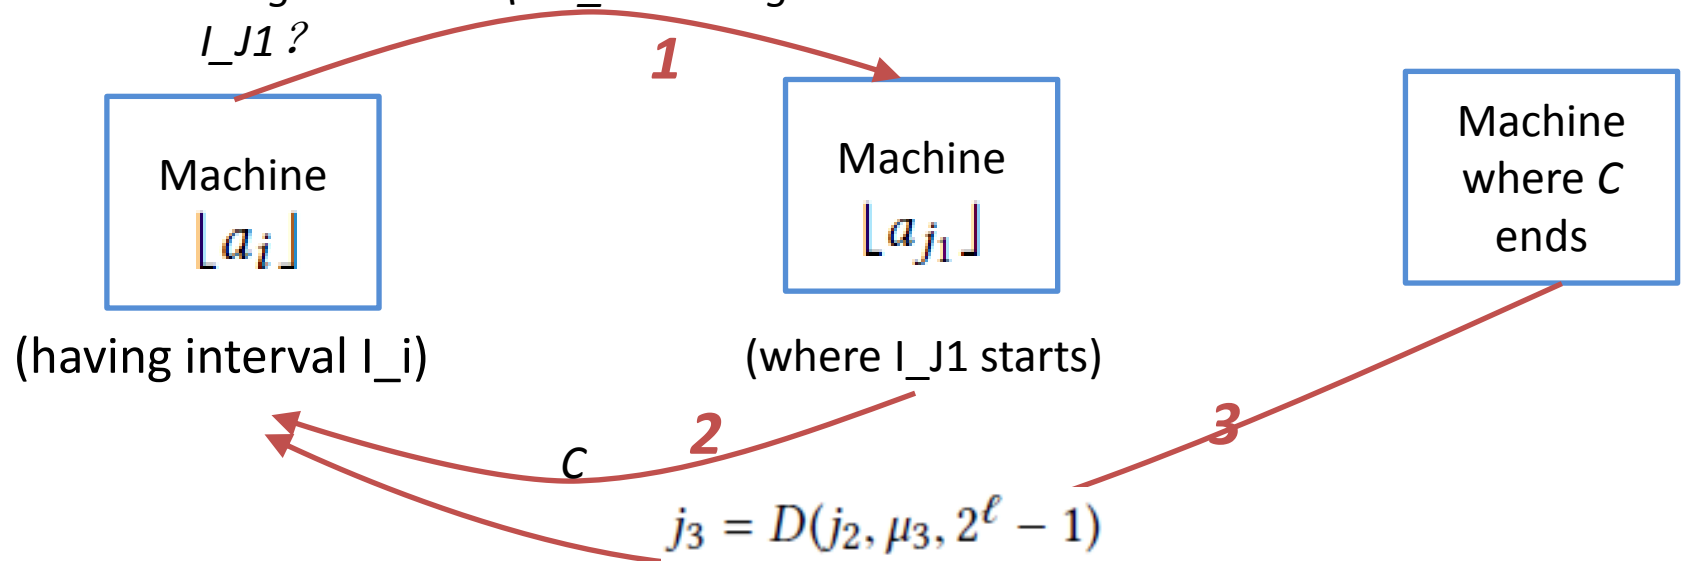


I_{j_2} is the earliest starting interval after the interval C ends,

Consider any triplet $\mu_1, \mu_2, \mu_3 \in W$ such that

$$(1 - \eta)\mu \leq \mu_1 + \mu_2 + \mu_3 \leq \mu \quad j_1 = D(i, \mu_1, 2^\ell - 1)$$

Which is the earliest ending crossing interval of weight at least μ_2 starting no earlier than I_{j_1} ?



I_{j_2} is the earliest starting interval after the interval C ends,

2, 3 requires only $O(1)$ rounds.

We take the minimum j_3 over all possible triplets in parallel.

There are at most

$|W|^3 = O(\log_{1+\eta}^3 n)$ triplets to be considered for each i . Thus, each machine uses memory at most $(n/m)O(\log_{1+\eta}^3 n) = \tilde{O}(n/m)$.

Prove the approximation ratio by induction.